

# Out of Memory in Containers

## Java Heap (JVM) vs. Native Heap (malloc)

Martijn Vermeulen



engineered  
for ambition

# Introduction

**adyen**

engineered  
for ambition

# About Adyen



# About me



# Containers & JVM memory challenges

# Upside of containers: Isolation

# Downside of containers: Isolation

**Upside of the JVM:  
Does a lot of memory management  
for you**

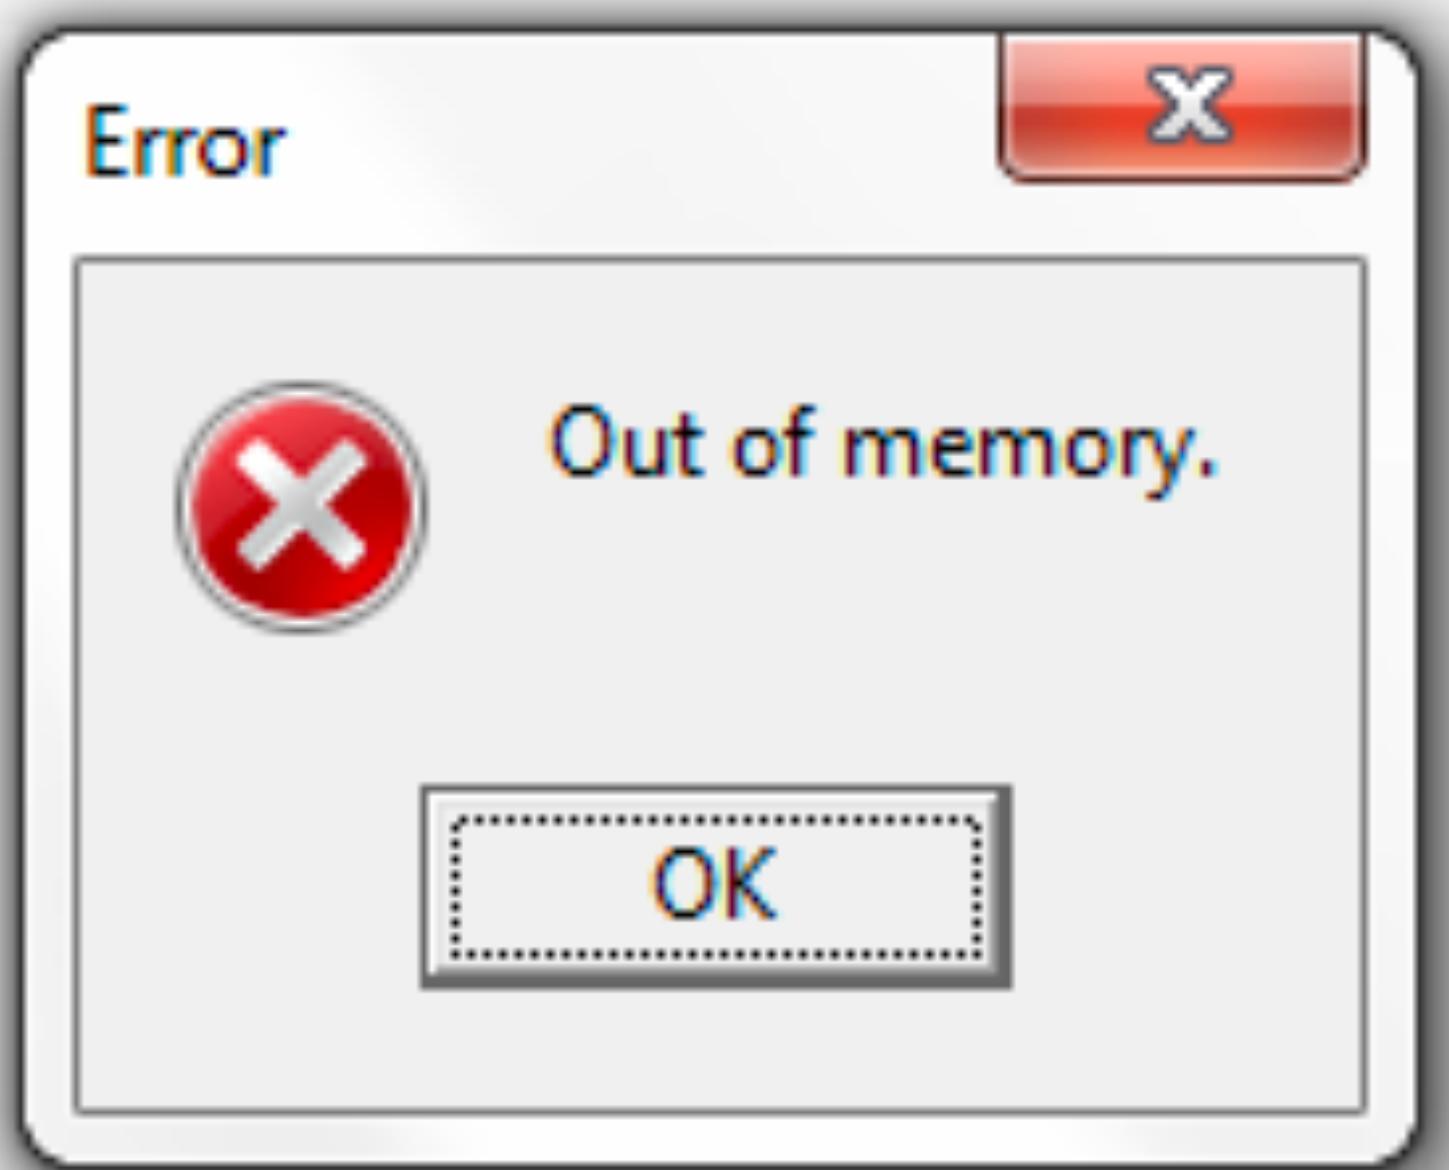
**Downside of the JVM:**  
**Does a lot of memory management**  
**for you**

# Investigating real issues

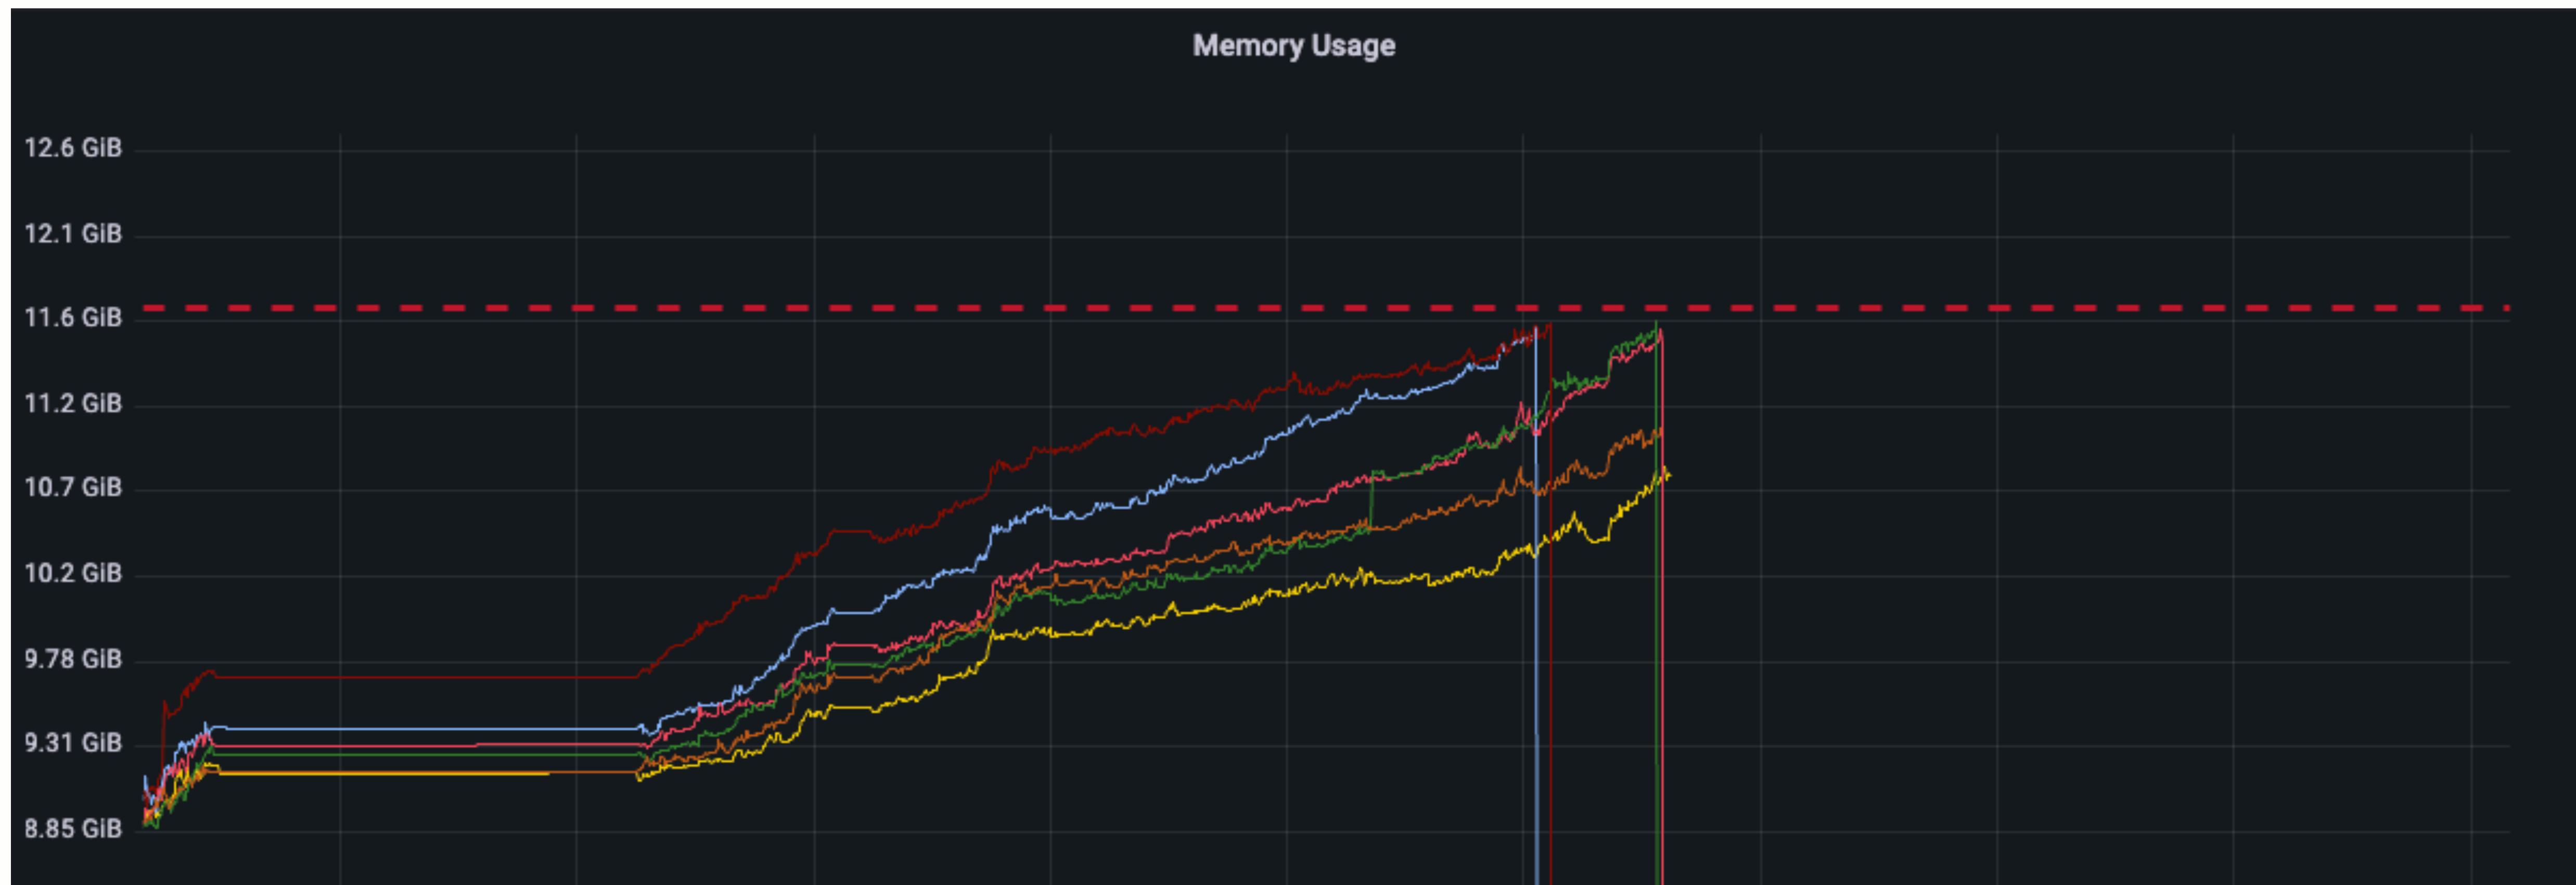
**adyen**

engineered  
for ambition

# The problem



# The memory usage

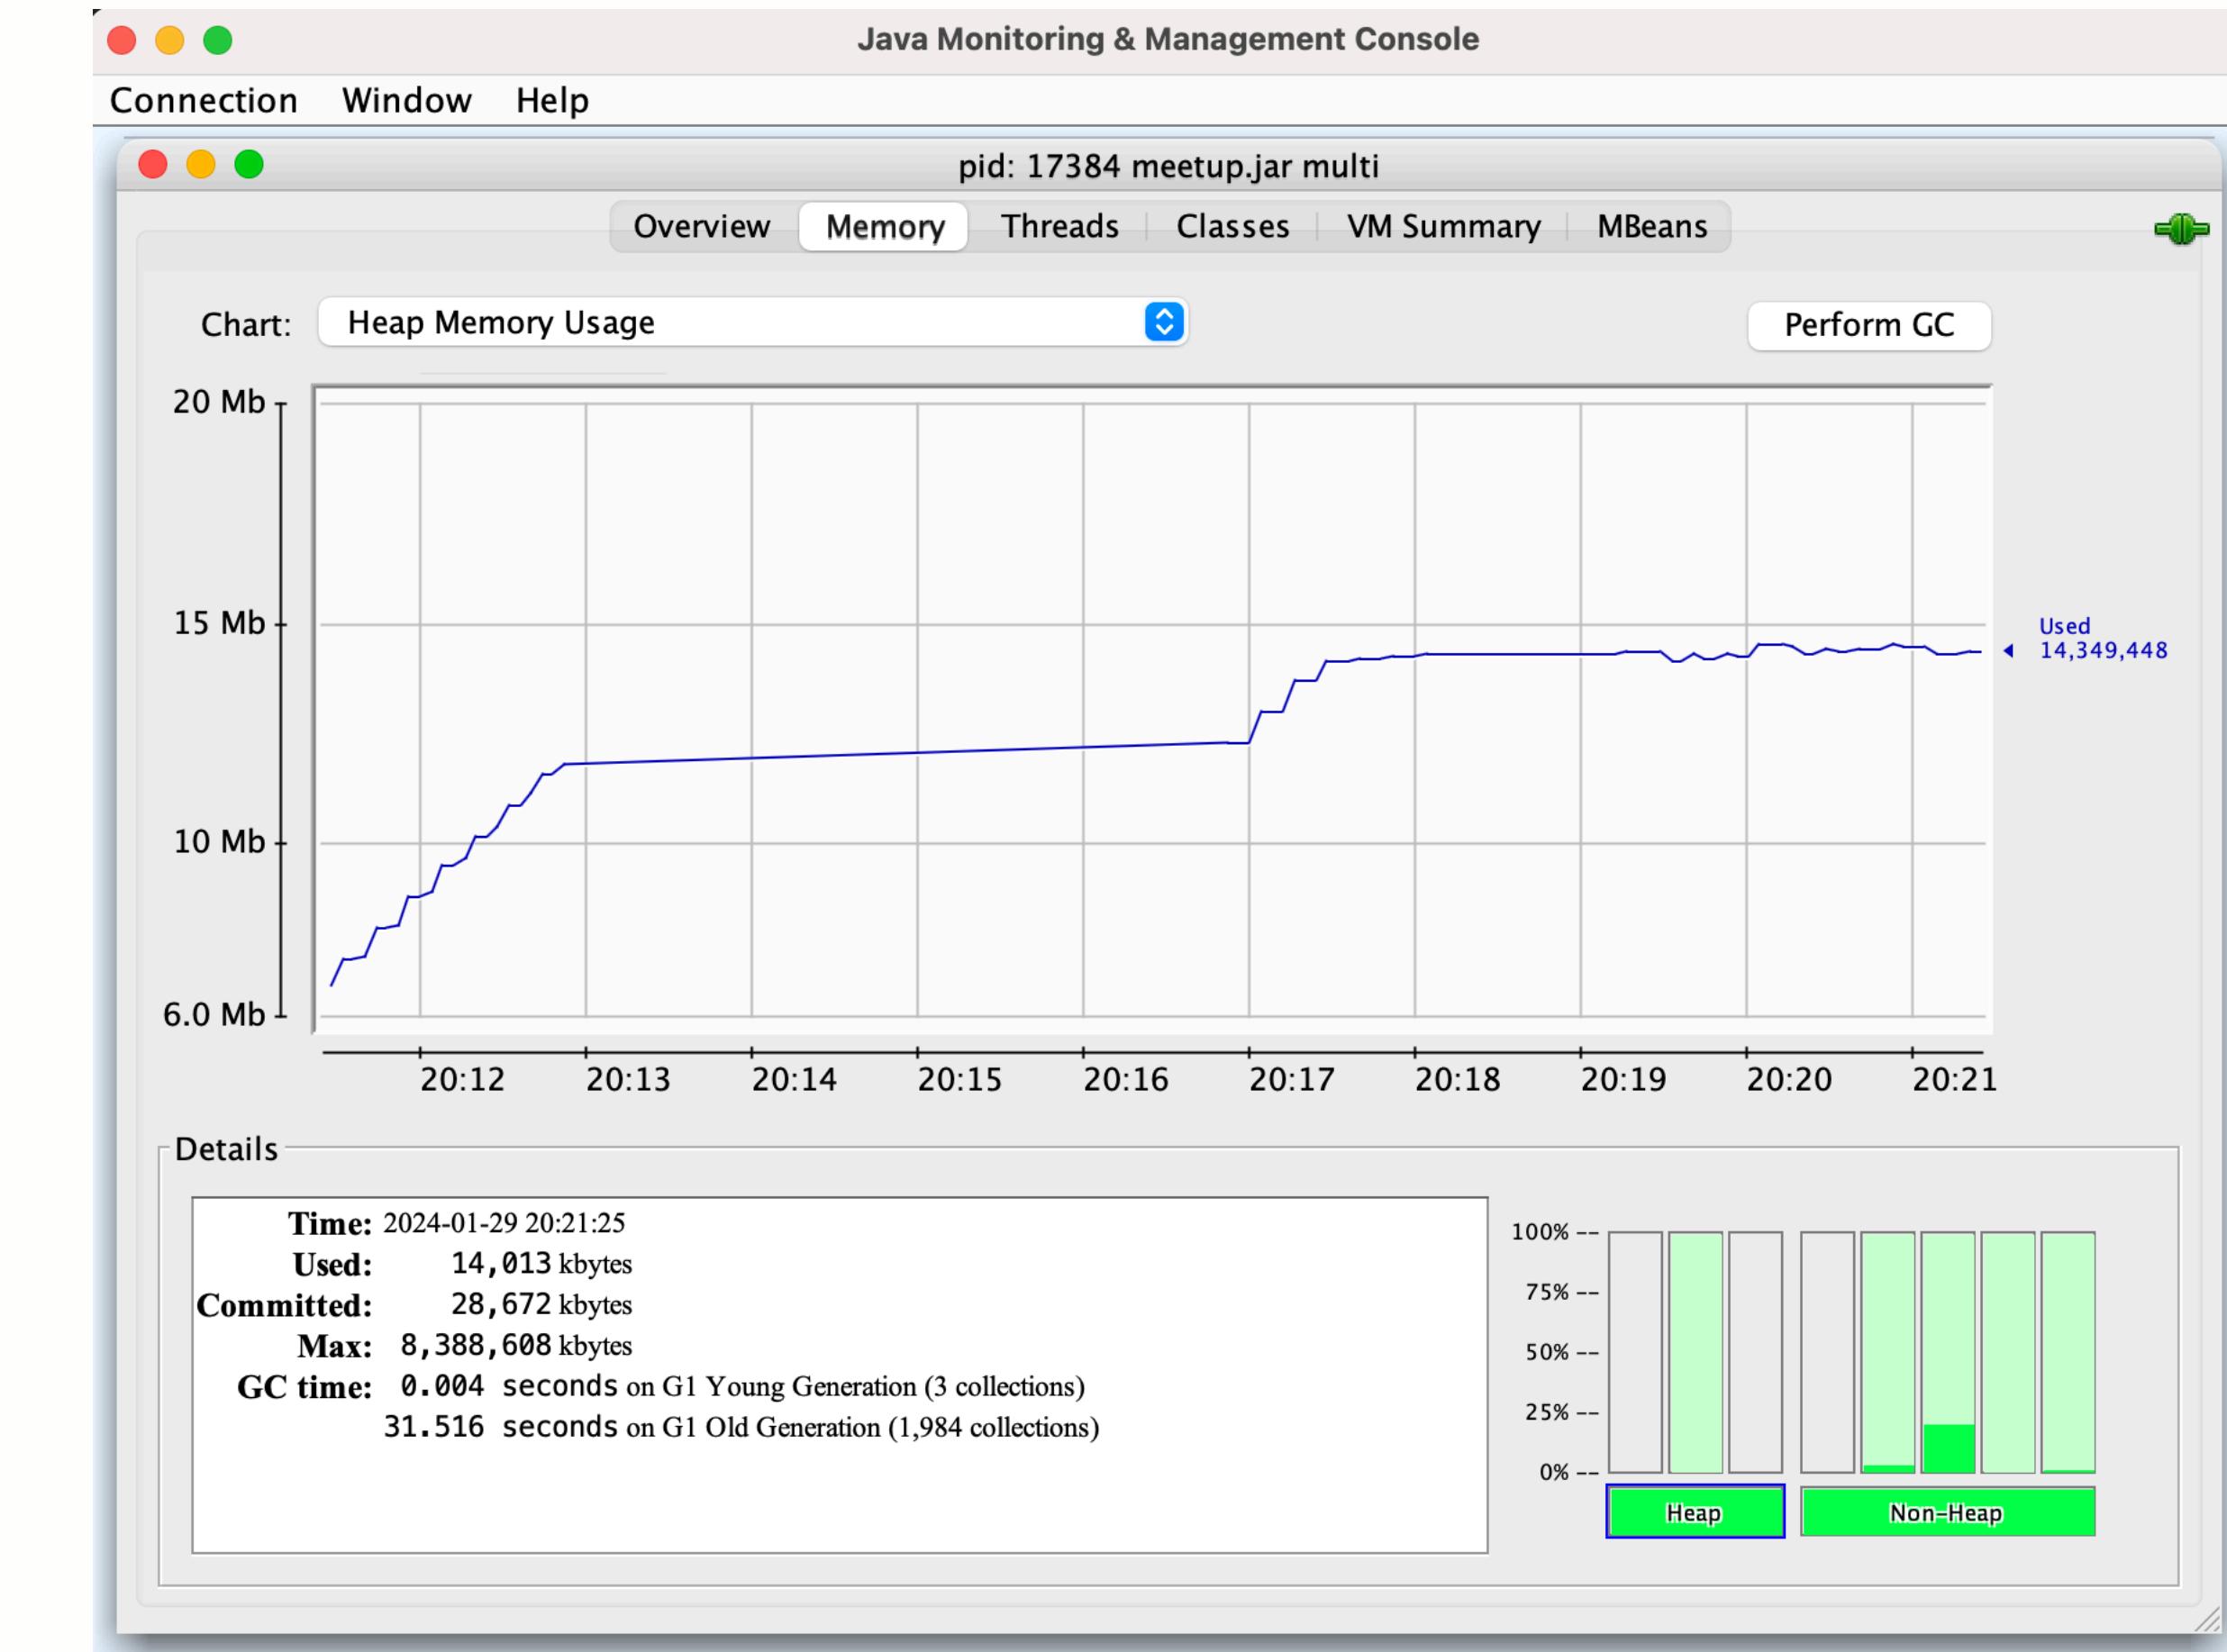
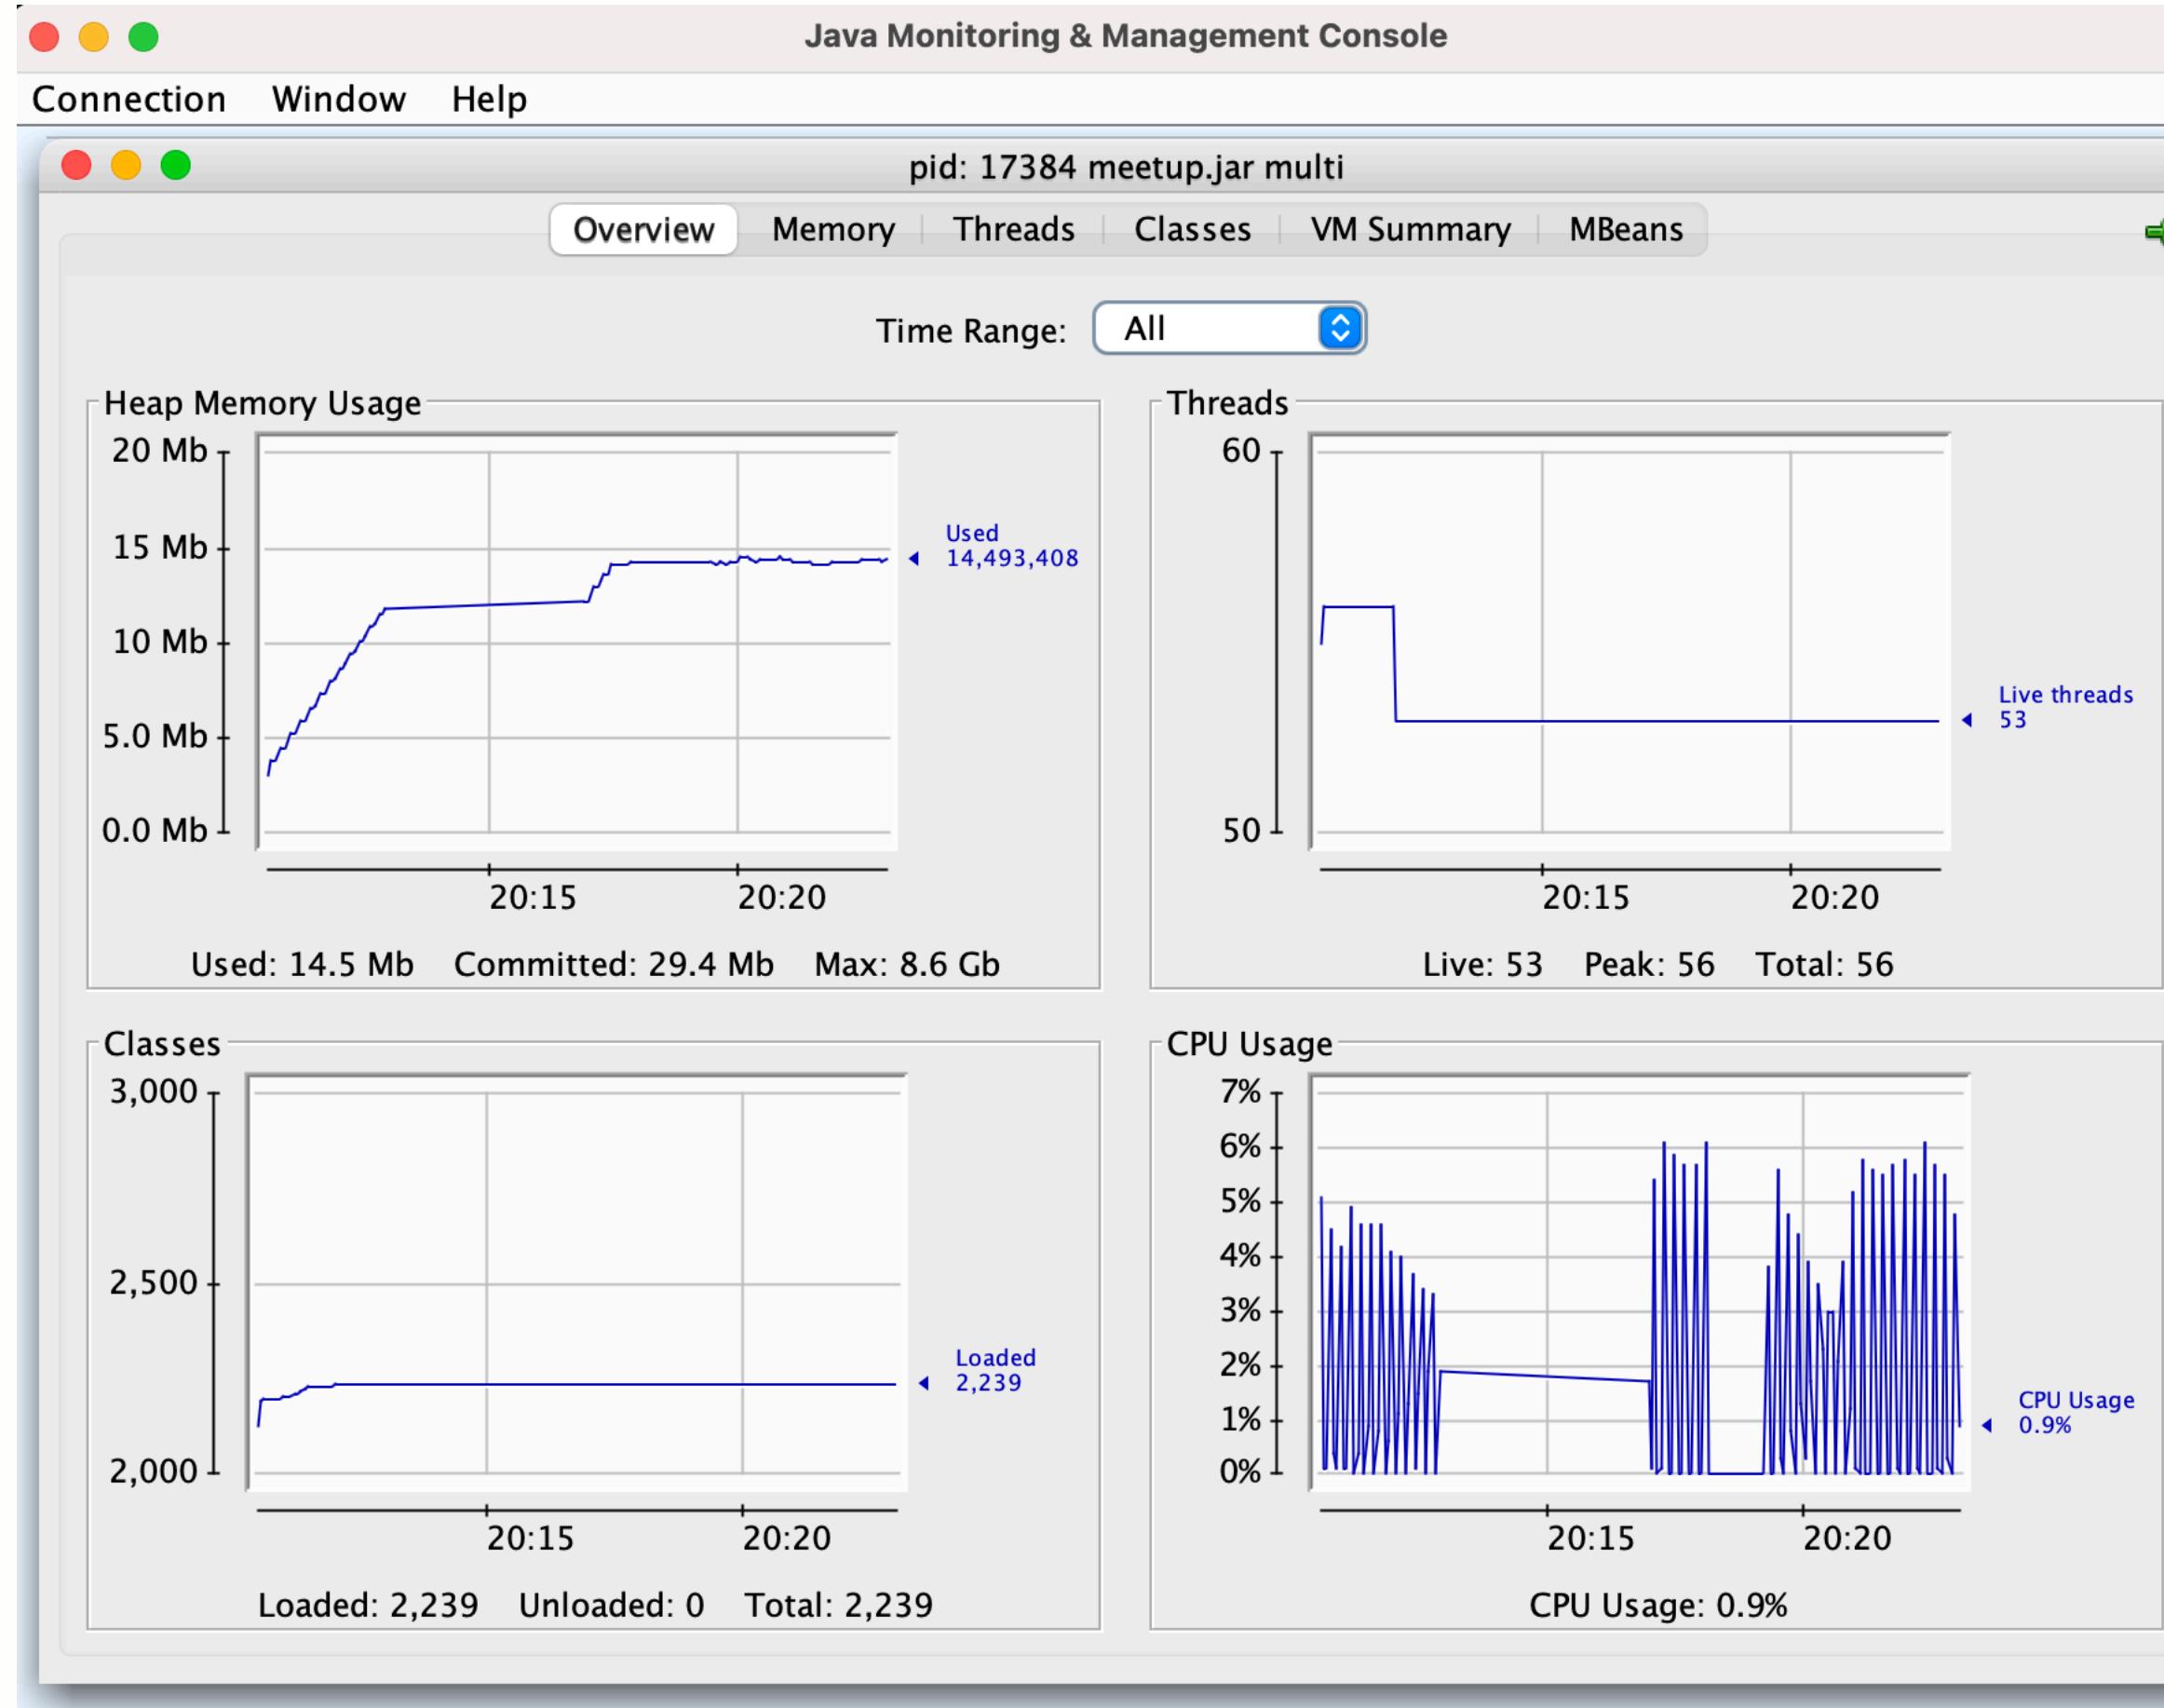


# The investigation

- Heap ?
- Threads ?
- Class ?
- Code cache?
- Other processes?

# The investigation

## JConsole



# The investigation

## jcmd, jstat, jmap & MAT

- **jcmd <pid> GC.heap\_info**
- **jstat -gc <pid>**
- **jmap -> Eclipse MAT**

# The investigation

# Native Memory Tracking

`summary` / `detail`

Detecting memory leaks  
with baseline -> diff

Overhead

# The investigation

# Native Memory Tracking

## Total: reserved, committed

- **malloc**
- **mmap**

- Java Heap
- Class
- Thread
- Code
- GC
- Compiler
- Internal
- Other
- Symbol
- Native Memory Tracking
- Shared class space
- Arena Chunk
- Module
- Safepoint
- Synchronization
- Serviceability
- Metaspace
- String Deduplication
- Object Monitors

# Native memory tracking

## Summary

```
sh-4.4$ jcmt 1 VM.native_memory summary  
1:
```

Native Memory Tracking:

(Omitting categories weighting less than 1KB)

Total: reserved=2487217KB, committed=81637KB

malloc: 13057KB #7182

mmap: reserved=2474160KB, committed=68580KB

- Java Heap (reserved=1048576KB, committed=1024KB  
(mmap: reserved=1048576KB, committed=1024KB)

- Class (reserved=1048688KB, committed=240KB  
(classes #842)  
( instance classes #720, array classes #122  
(malloc=112KB #1317)  
(mmap: reserved=1048576KB, committed=1024KB  
( Metadata: )  
( reserved=8192KB, committed=704KB  
( used=624KB)  
( waste=80KB =11.38%)  
( Class space: )  
( reserved=1048576KB, committed=1024KB  
( used=36KB)  
( waste=92KB =71.73%)

# Native memory tracking

## Summary

Shared class space (reserved=12288KB, committed=12116KB)  
(mmap: reserved=12288KB, committed=12116KB)

Arena Chunk (reserved=2329KB, committed=2329KB)  
(malloc=2329KB)

Module (reserved=14KB, committed=14KB)  
(malloc=14KB #402)

Safepoint (reserved=8KB, committed=8KB)  
(mmap: reserved=8KB, committed=8KB)

Synchronization (reserved=35KB, committed=35KB)  
(malloc=35KB #474)

Serviceability (reserved=1KB, committed=1KB)  
(malloc=1KB #6)

Metaspace (reserved=8203KB, committed=715KB)  
(malloc=11KB #9)  
(mmap: reserved=8192KB, committed=704KB)

String Deduplication (reserved=1KB, committed=1KB)  
(malloc=1KB #8)

Object Monitors (reserved=2KB, committed=2KB)  
(malloc=2KB #8)

# Native memory tracking

Baseline

vs.

Diff (compared to the baseline)

```
sh-4.4$ jccmd 1 VM.native_memory baseline  
1:  
Baseline succeeded  
sh-4.4$ jccmd 1 VM.native_memory summary.diff  
1:
```

Native Memory Tracking:

(Omitting categories weighting less than 1KB)

Total: reserved=2469091KB +675KB, committed=76171KB +675KB

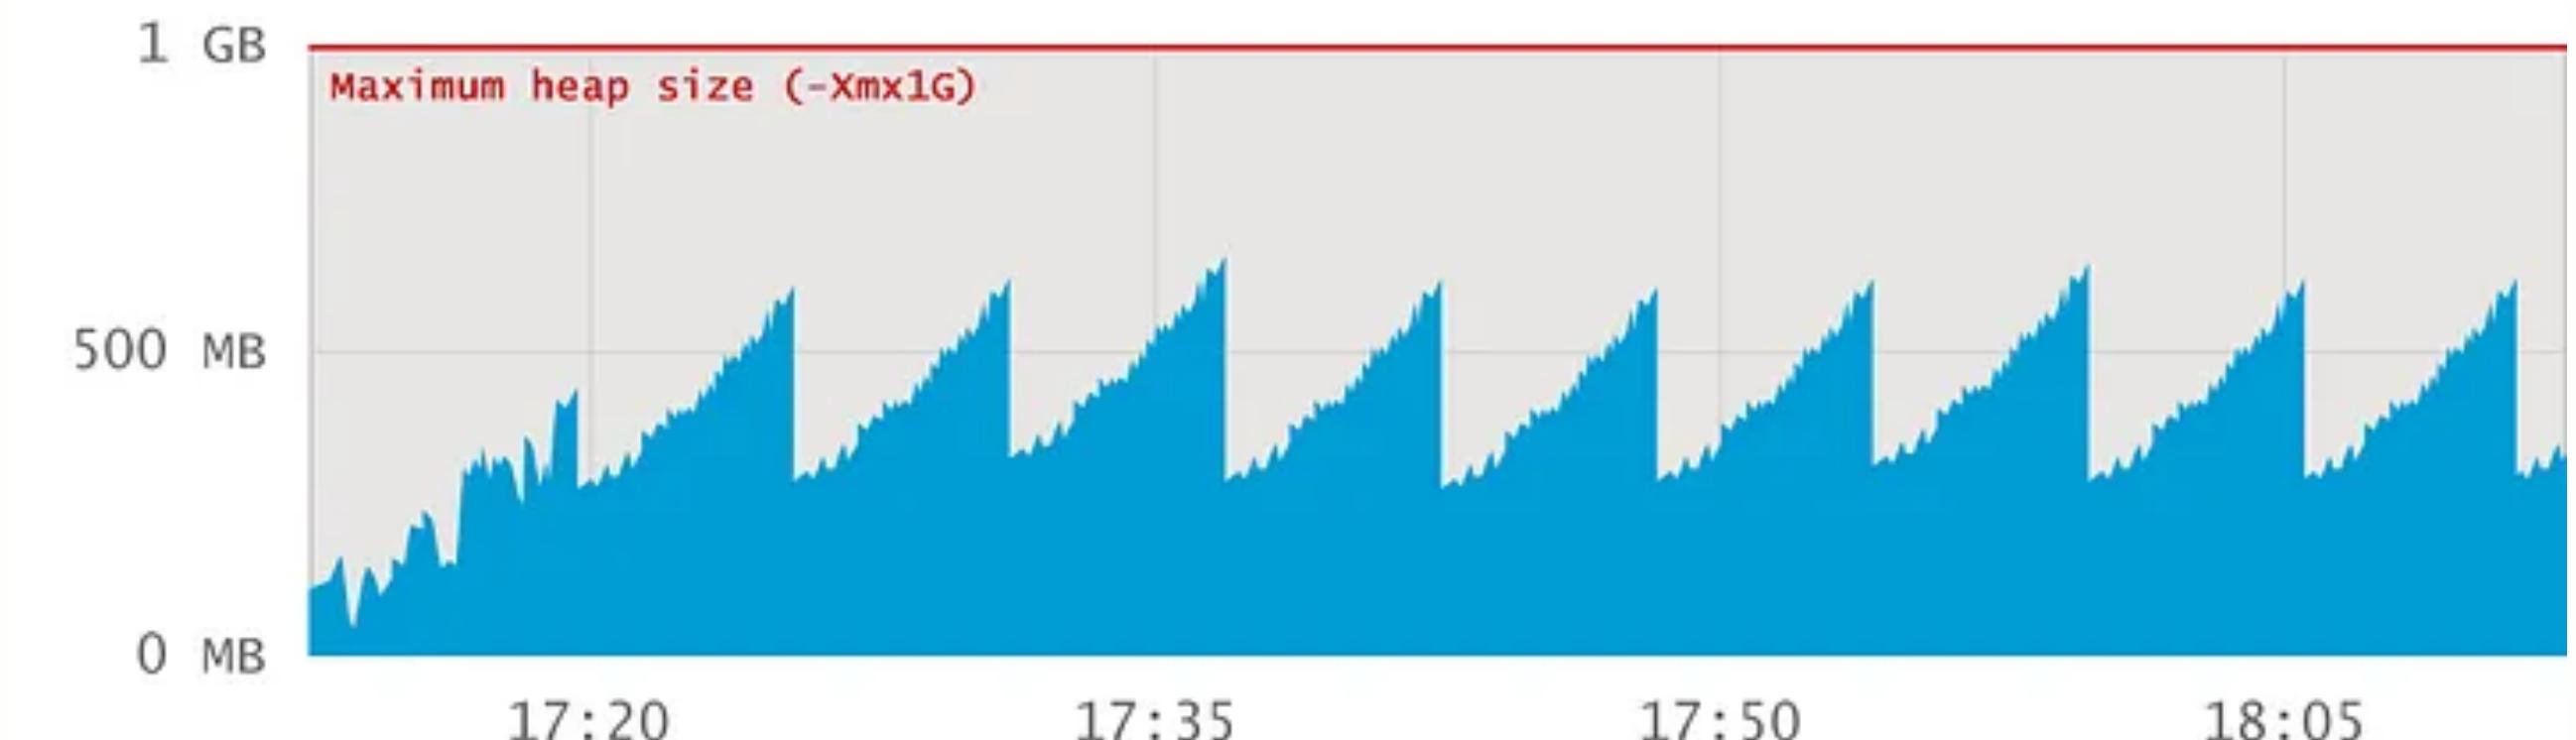
- Java Heap (reserved=1048576KB, committed=1048576KB  
(mmap: reserved=1048576KB, committed=1048576KB)

- Class (reserved=1048676KB, committed=1048676KB  
(classes #791)  
( instance classes #673, array  
(malloc=100KB #1086 +4)  
(mmap: reserved=1048576KB, committed=1048576KB)

# The investigation JVM related memory usage

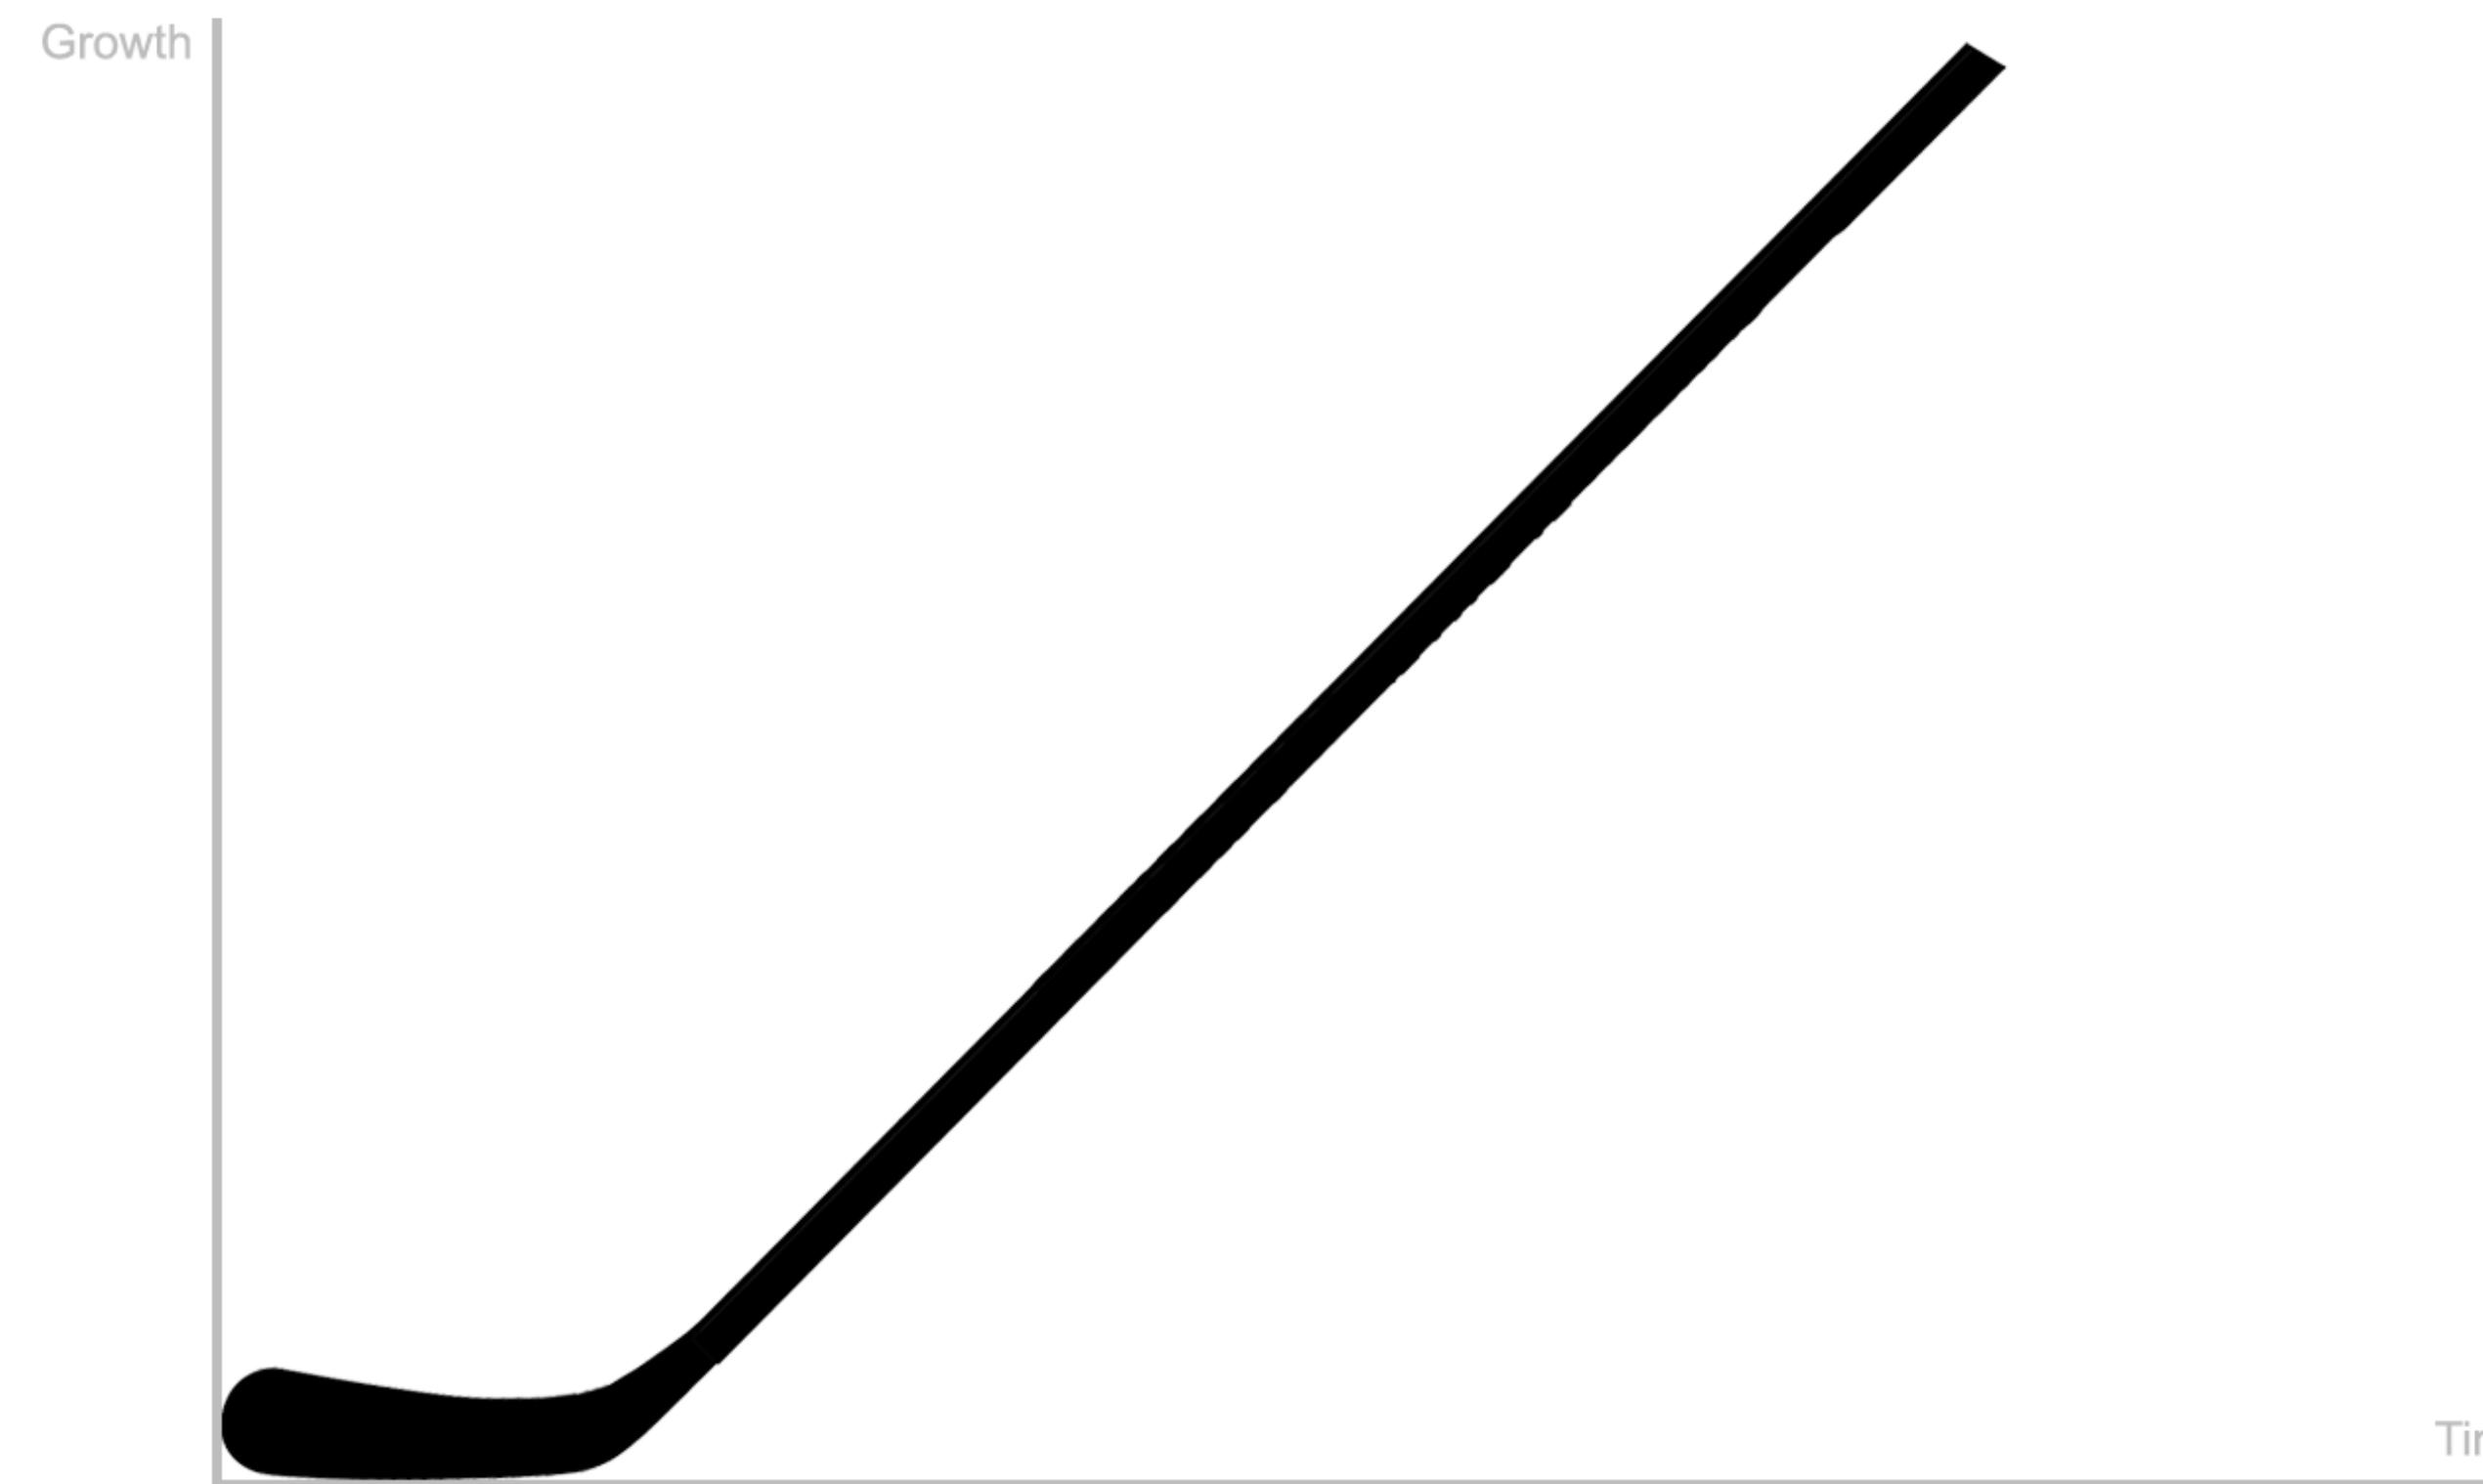


<https://medium.com/medical-myths-and-models/can-you-shock-a-heart-that-has-flatlined-d156a12a898c>



<https://medium.com/@plumbr/memory-leaks-fallacies-and-misconce-8c79594a3986>

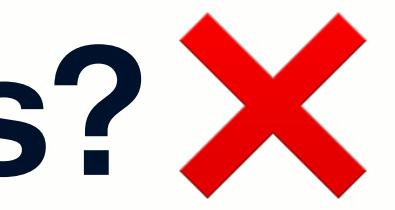
# The investigation top & docker stats



# The investigation

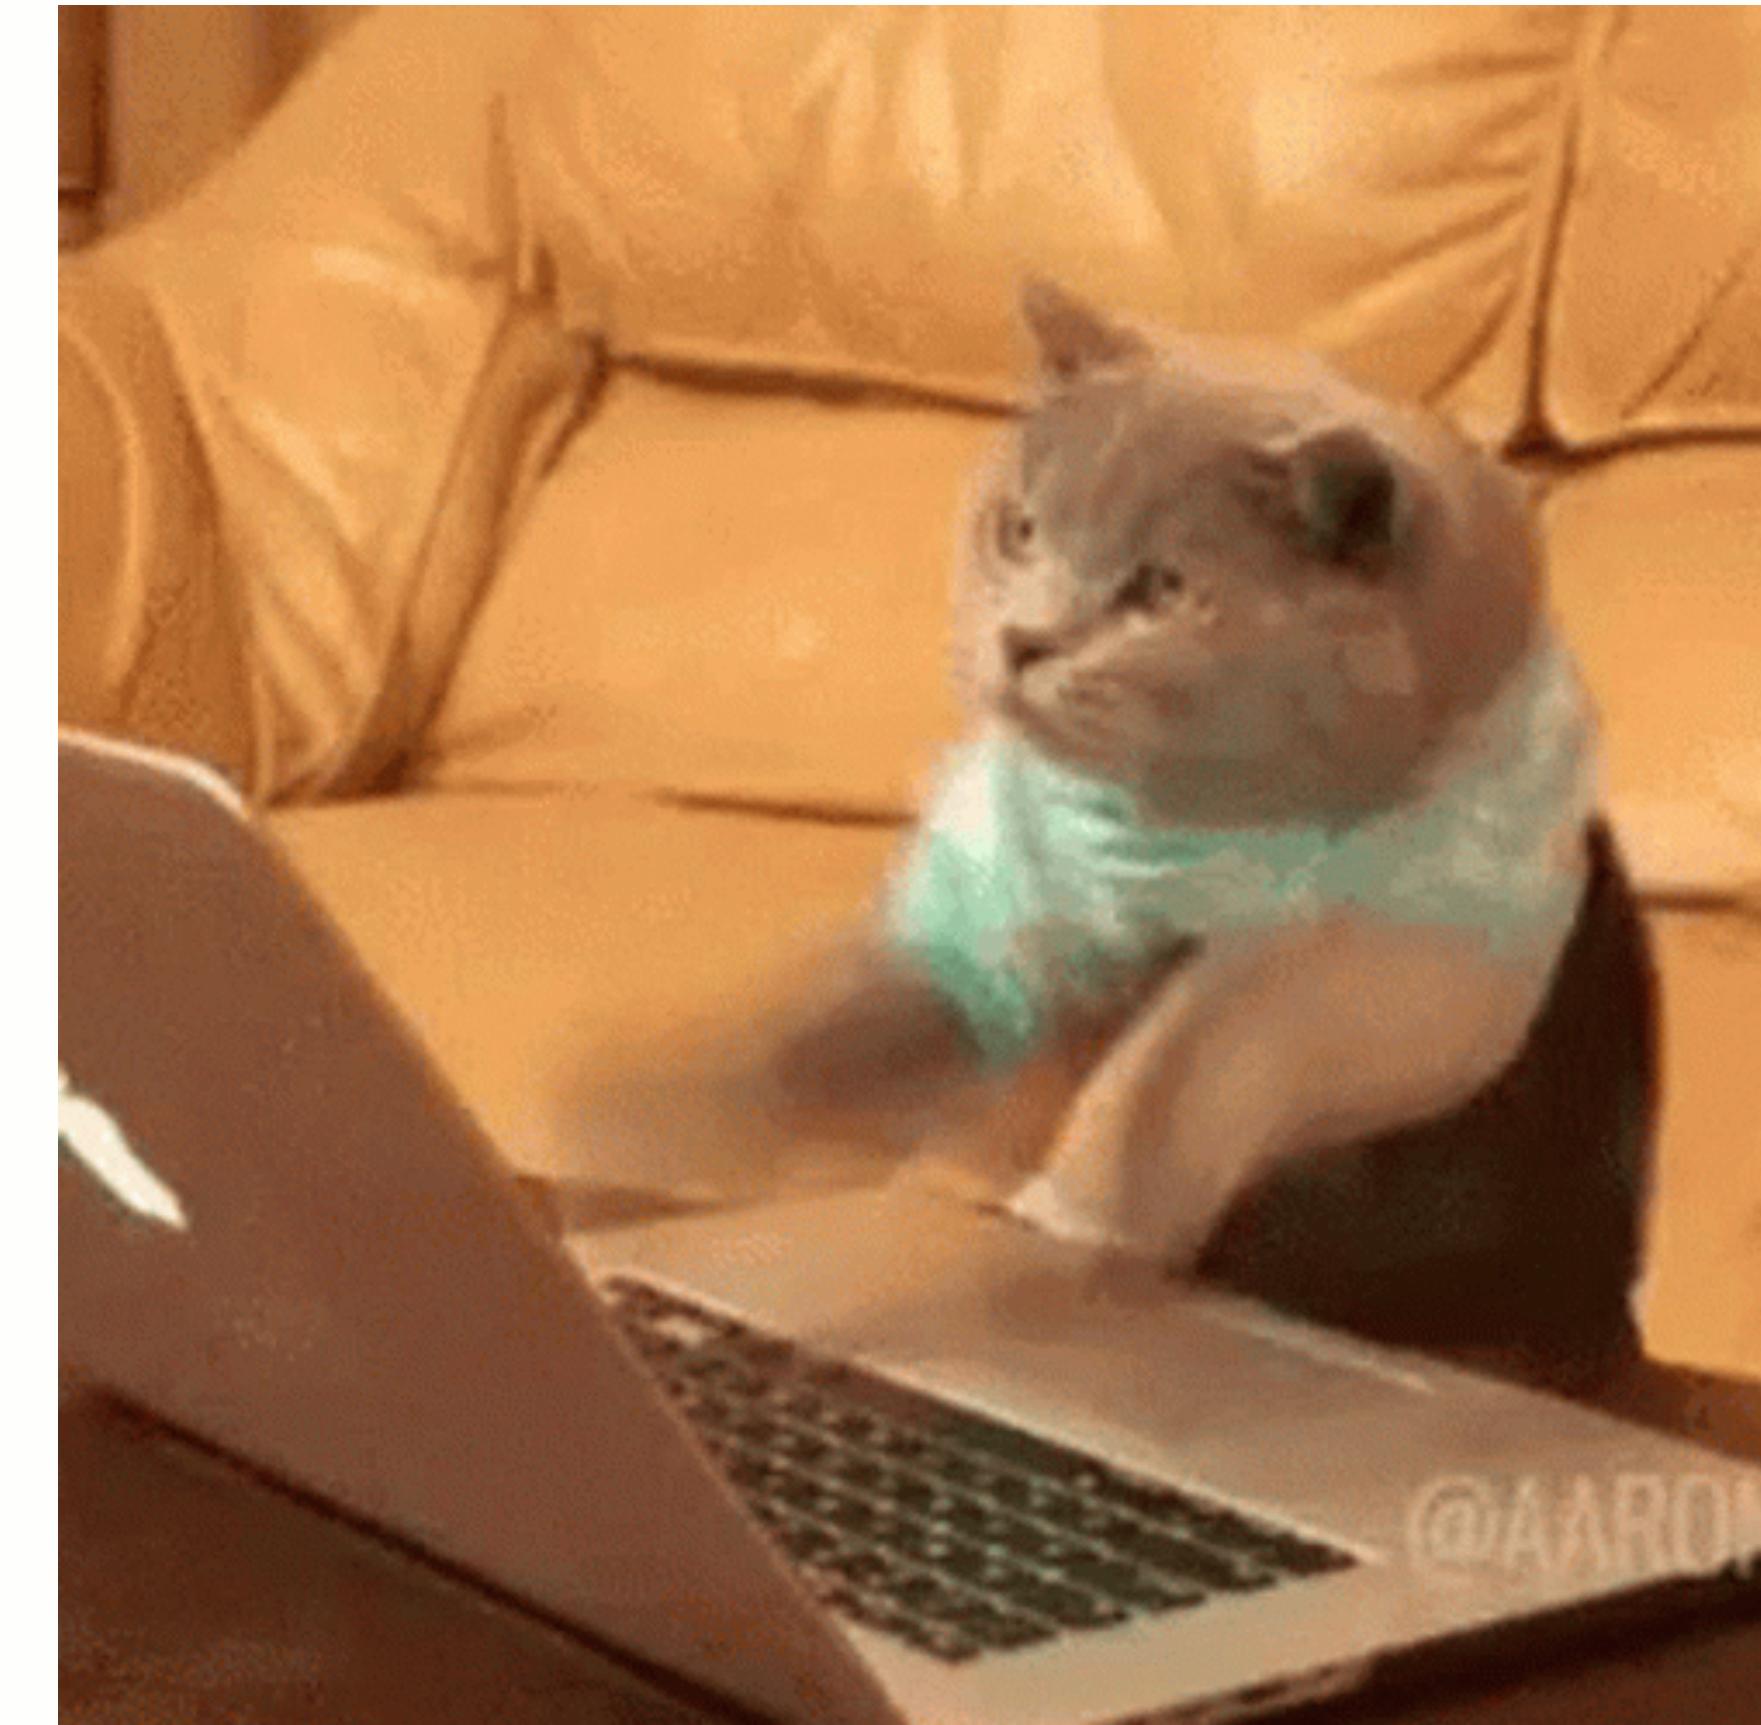
- Heap ? 
- Threads ? 
- Class ? 
- Code cache? 
- Other processes?

# The investigation

- Heap ? 
- Threads ? 
- Class ? 
- Code cache? 
- Other processes? 

# The culprit?

- Trying stuff...
- Google

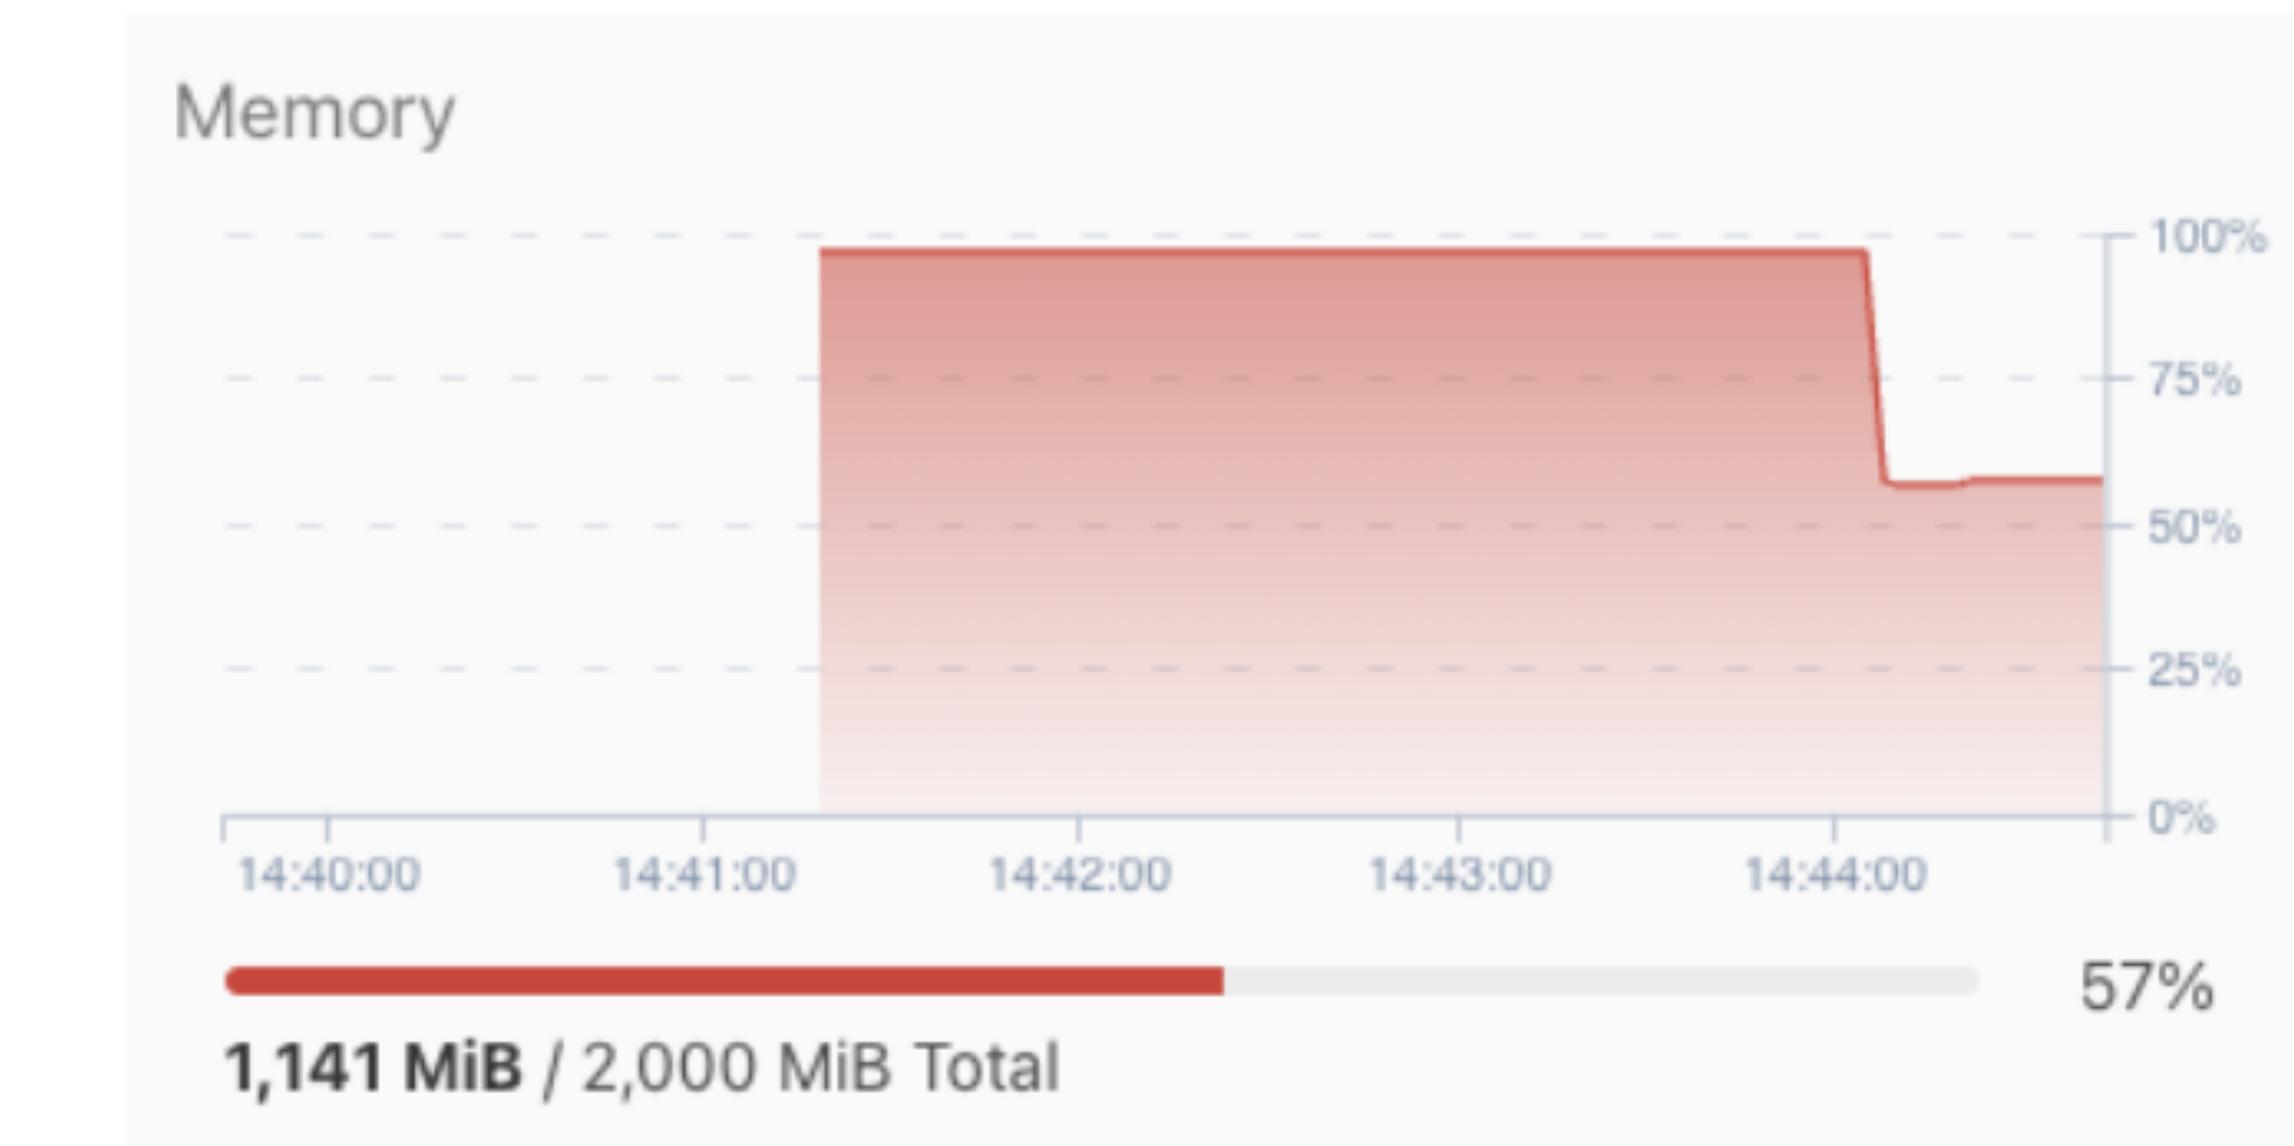


<https://tenor.com/view/typing-cat-busy-writing-chatting-gif-15923952>

# The culprit!

- **System.trim\_native\_heap?**
- **Native heap overhead!** 

```
bash-4.4$ jcmb 1 System.trim_native_heap
1:
Attempting trim...
Done.
Virtual size before: 10757268k, after: 10757140k, (-128k)
RSS before: 2009392k, after: 1191544k, (-817848k)
Swap before: 0k, after: 0k, (0k)
```



# The culprit!

- Native heap overhead?



Surprised Engineer Stock Image  
<https://www.shutterstock.com/search/surprised-engineer>

# Native heap / malloc

**“But Martijn, we are at a Java Meetup?”**

I know!

# JVM & Native heap

## Java heap

- Java heap is allocated using mmap (mostly)
  - Max size is preallocated as one contiguous area
- JVM manages this memory directly

# JVM & Native heap

## Native heap

- The C way
- Allocated by using the underlying malloc and free

# JVM & Native heap

## Native heap

- GUI
- Buffers for data compression routines
- Compiled code generated by JIT Compiler
- Threads to map to Java threads
- Malloc allocations by application JNI code

# JVM & Native heap

## Malloc in JNI code

- `java.nio.DirectByteBuffer`
- `sun.misc.Unsafe`
- Any native code potentially

# malloc & free

- **malloc**
  - **allocate x bytes of memory and return address**
- **free**
  - **frees the memory space at given address**

# malloc & free

- **malloc(100kb) -> addr1**
- **malloc(40kb) -> addr2**
- **malloc(1kb) -> addr3**

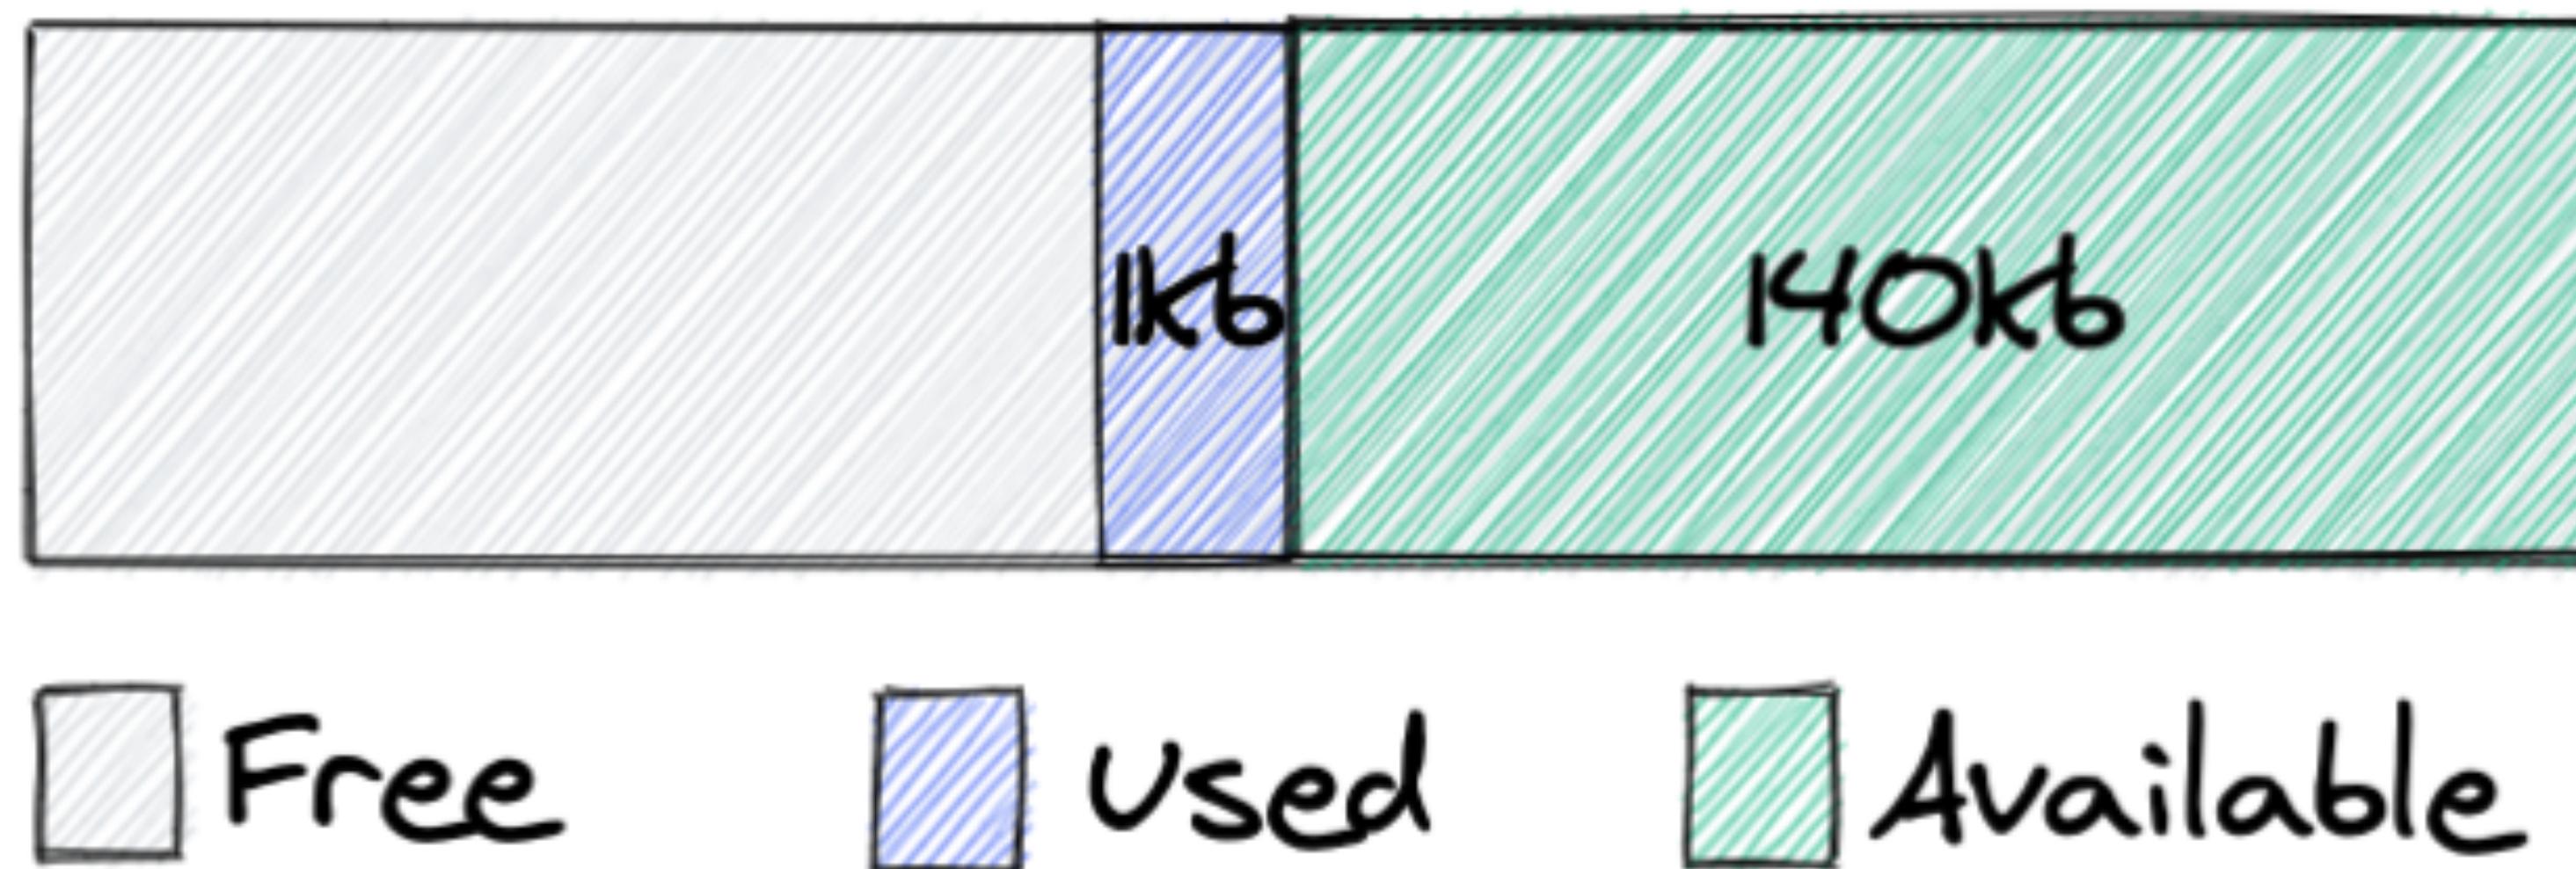


Free

Used

# malloc & free

- **free(addr1)**
- **free(addr2)**



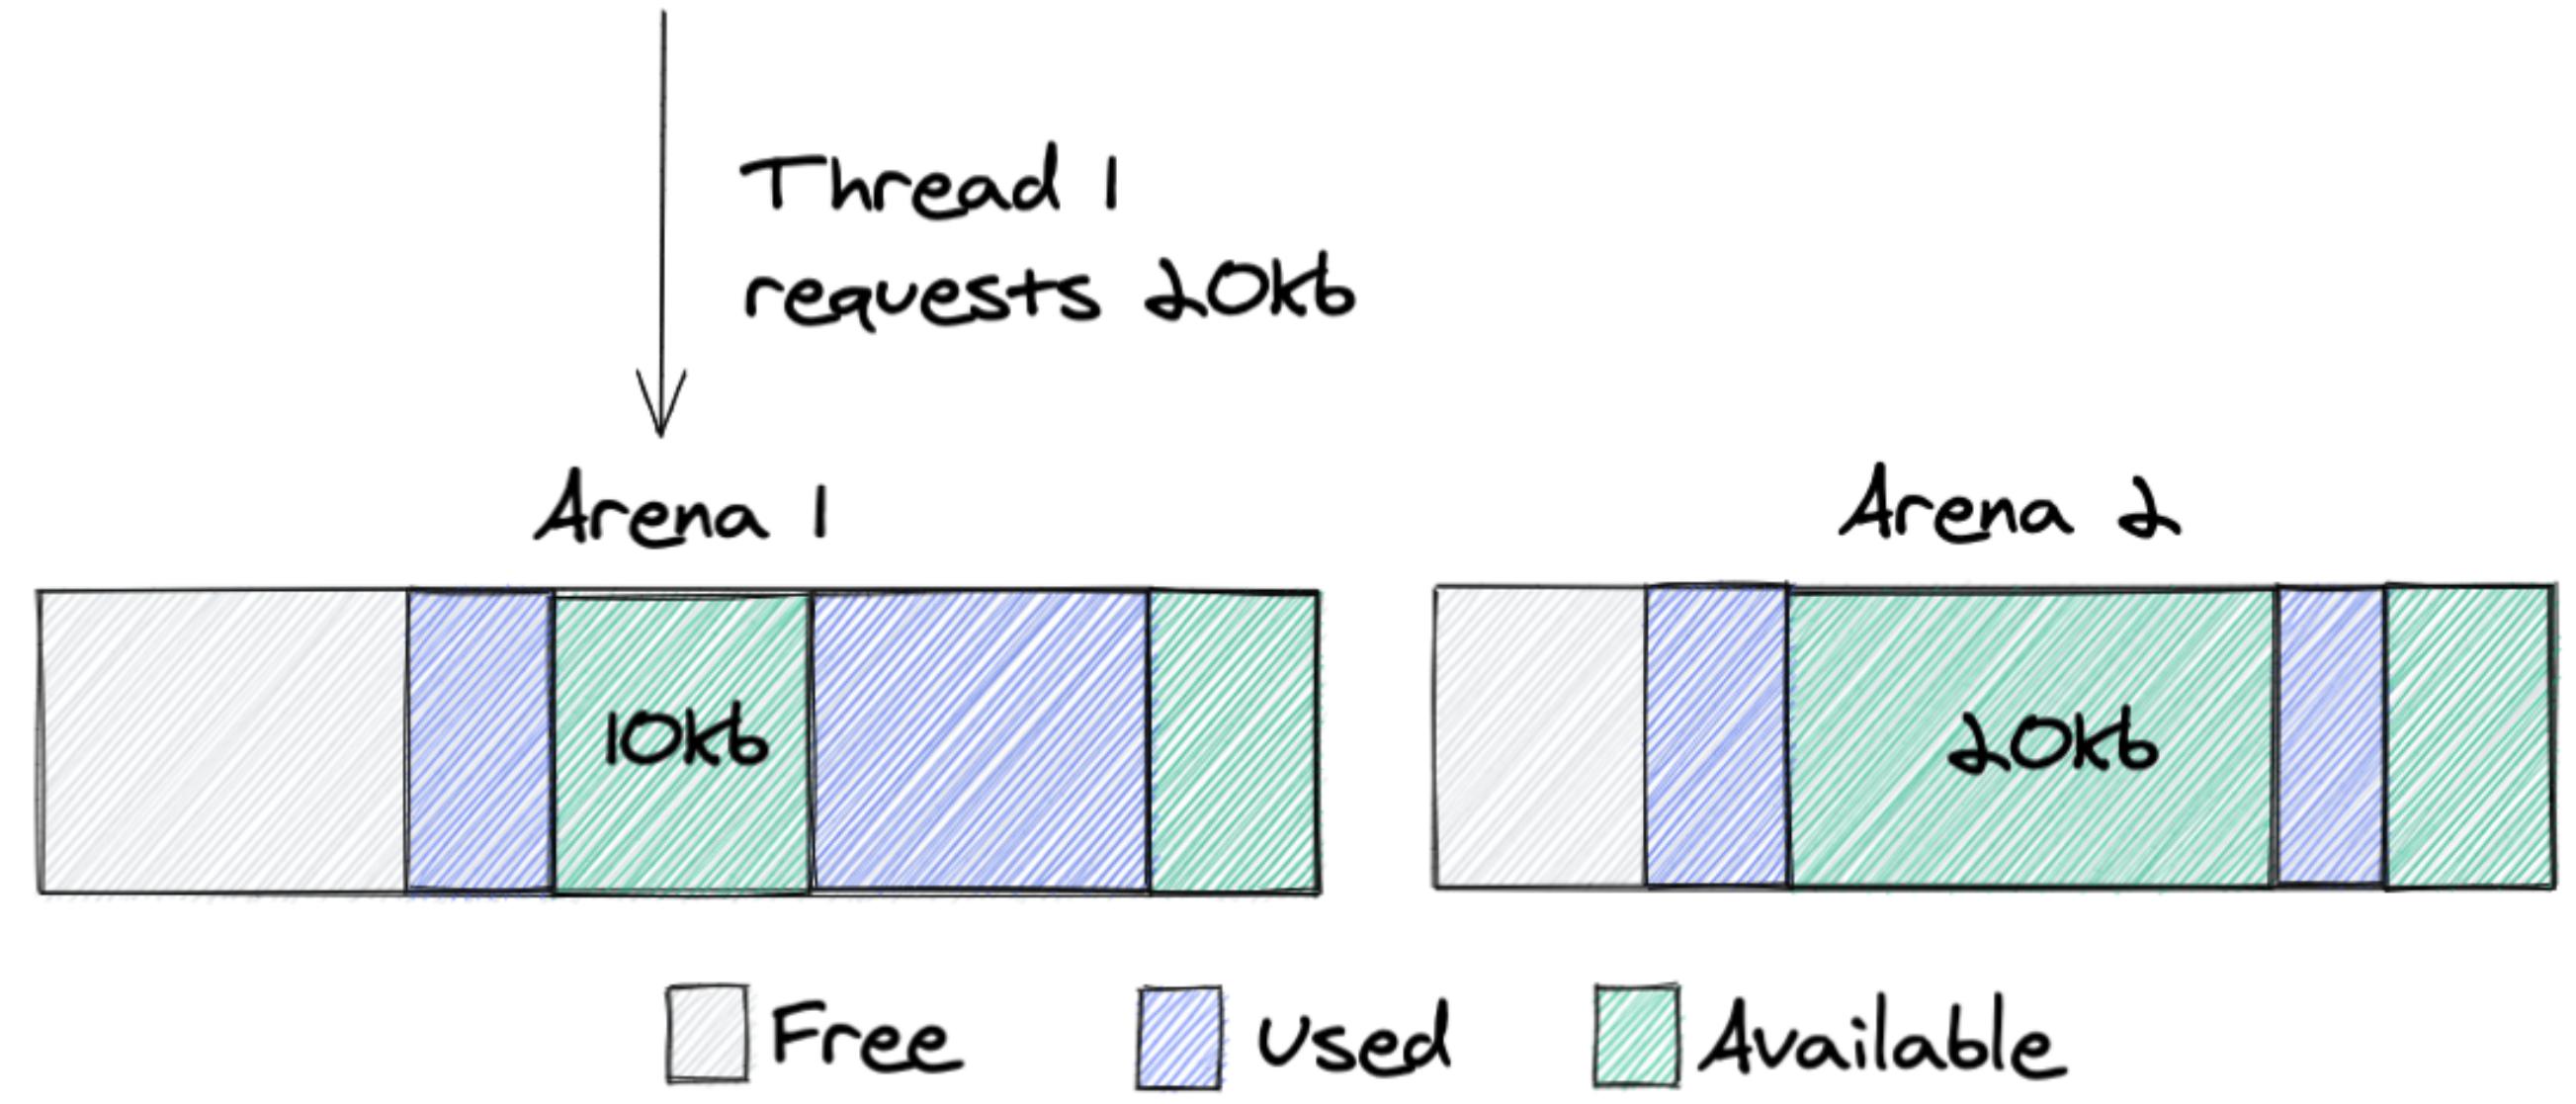
# glibc

- **GNU Project's implementation of the C standard library**
- **Default implementation of many Linux distros**

# glibc malloc

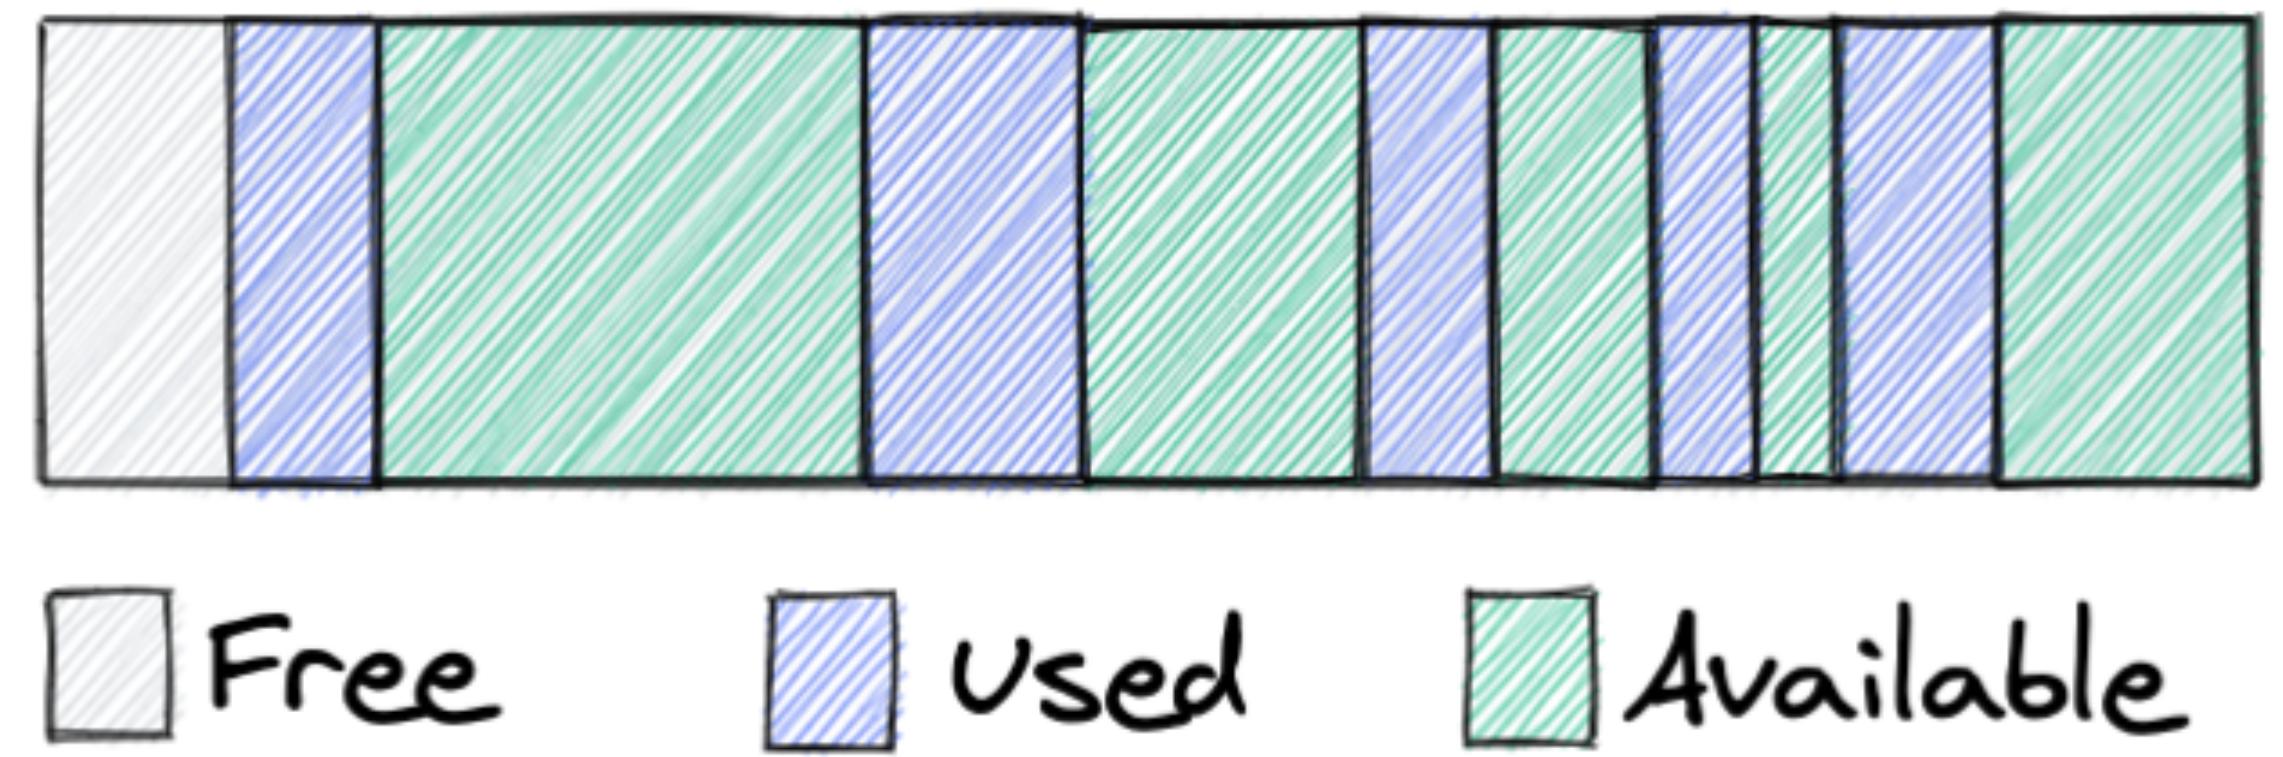
- **Uses multiple ‘arena’s to reduce locking**
- **One arena can contain multiple ‘heaps’ underneath**
- **By default, maximum of  $8 * \# \text{ cores}$  arenas**

# glibc malloc

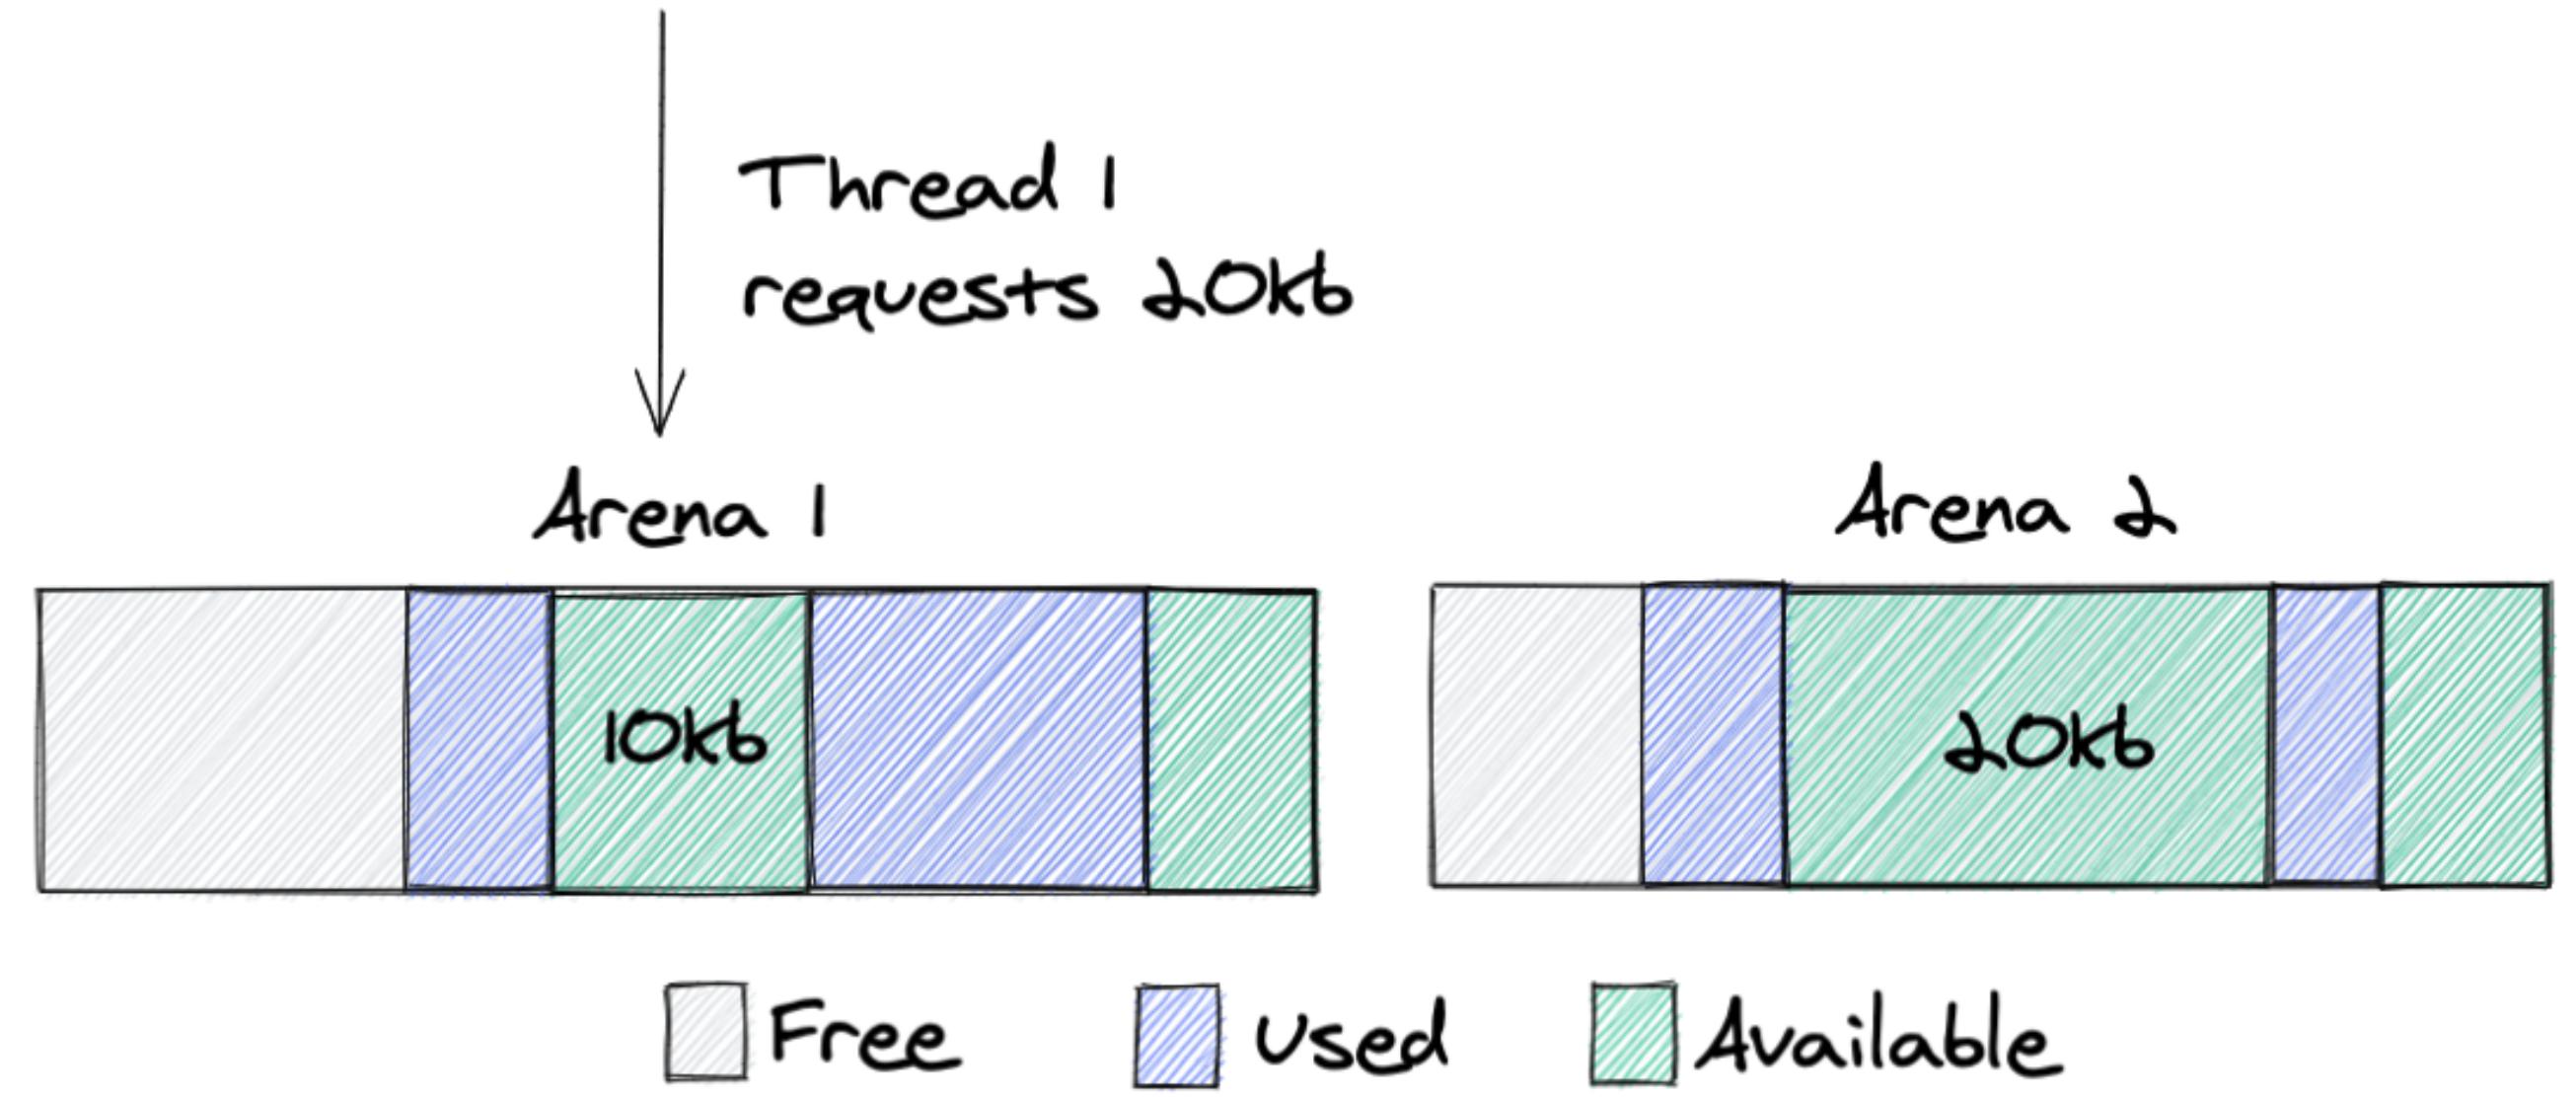


<https://blog.cloudflare.com/the-effect-of-switching-to-tcmalloc-on-rocksdb-memory-use>

# Memory fragmentation



# glibc malloc

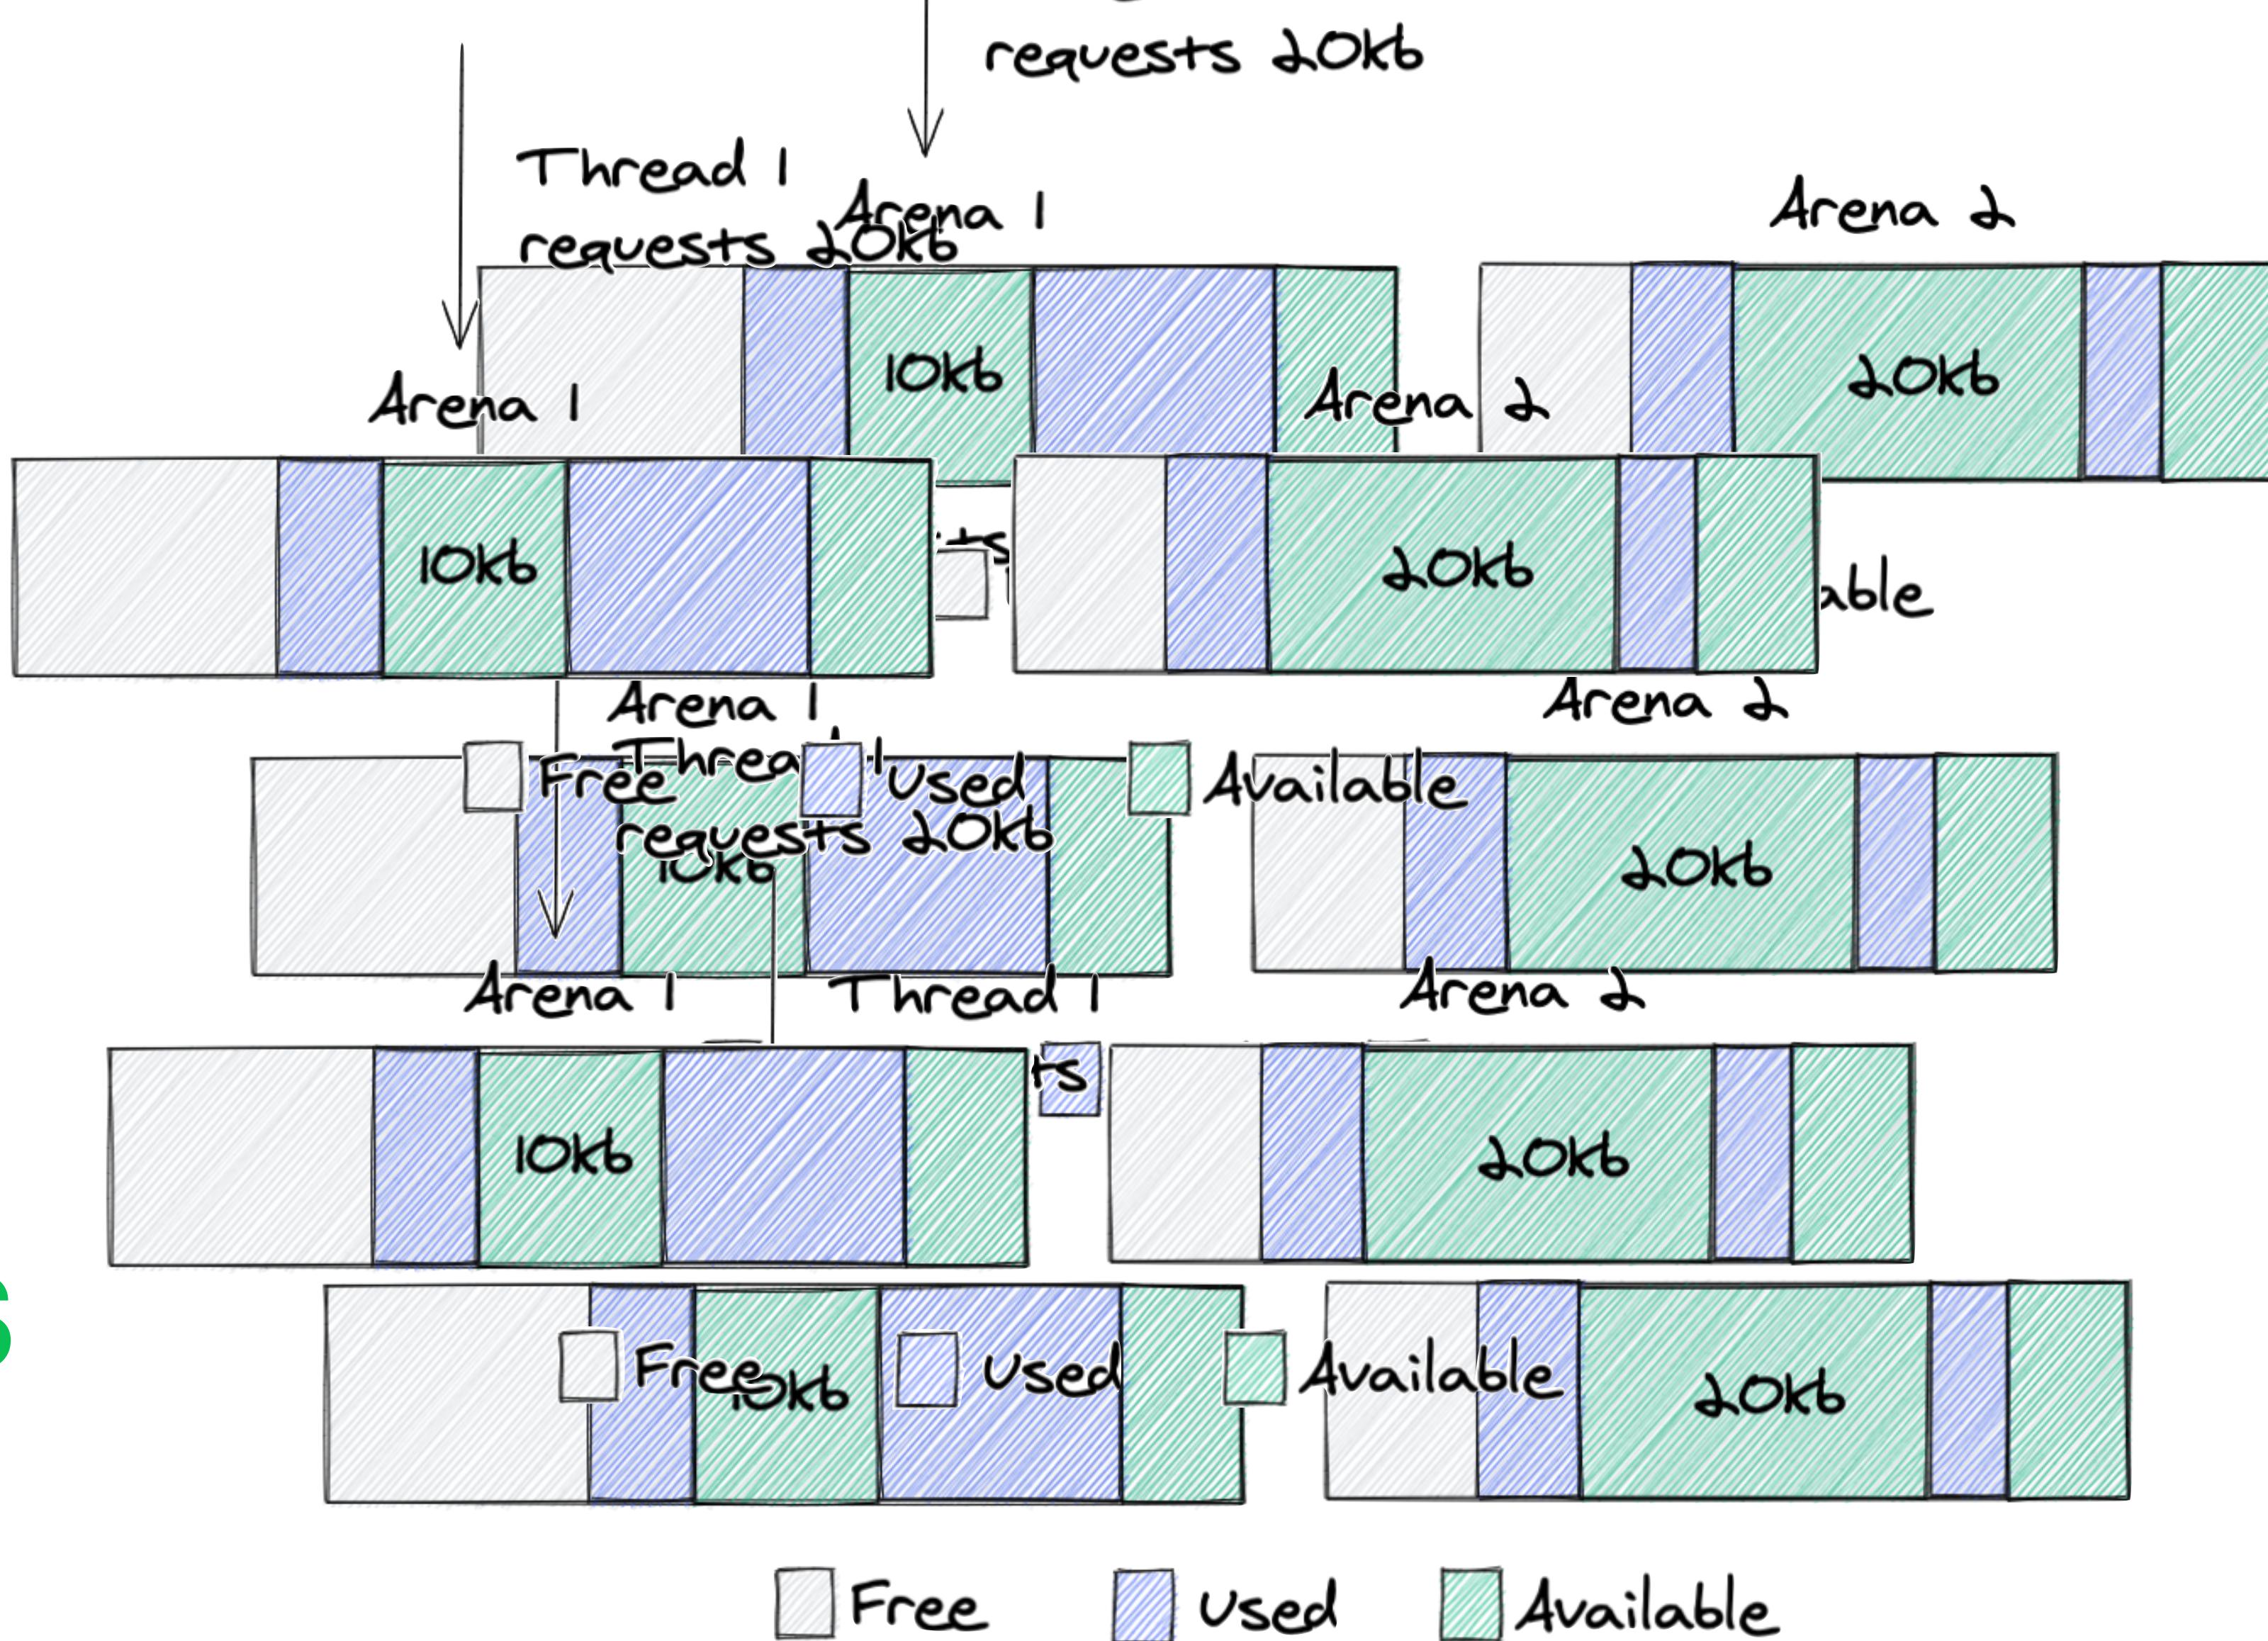


<https://blog.cloudflare.com/the-effect-of-switching-to-tcmalloc-on-rocksdb-memory-use>

# glibc malloc

+

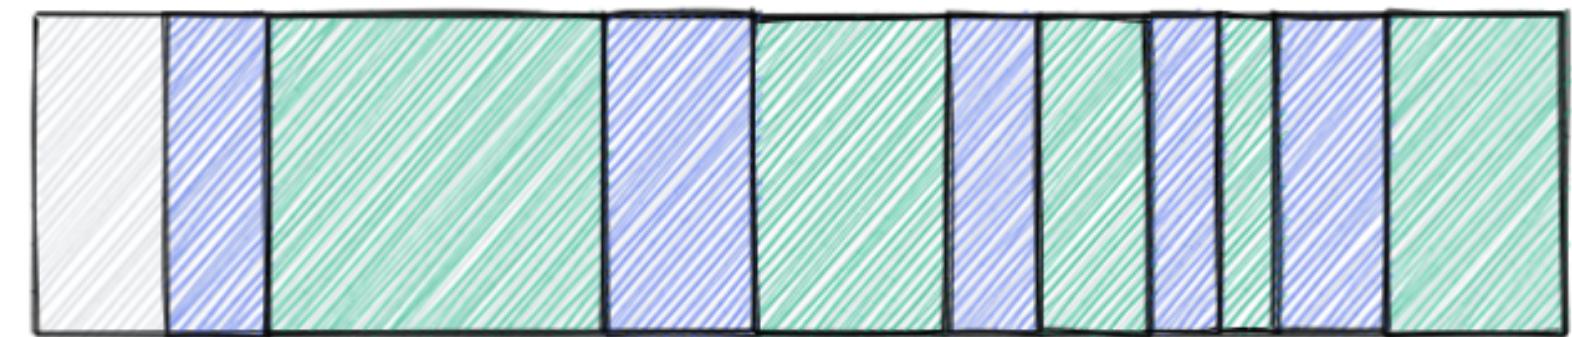
# Many threads



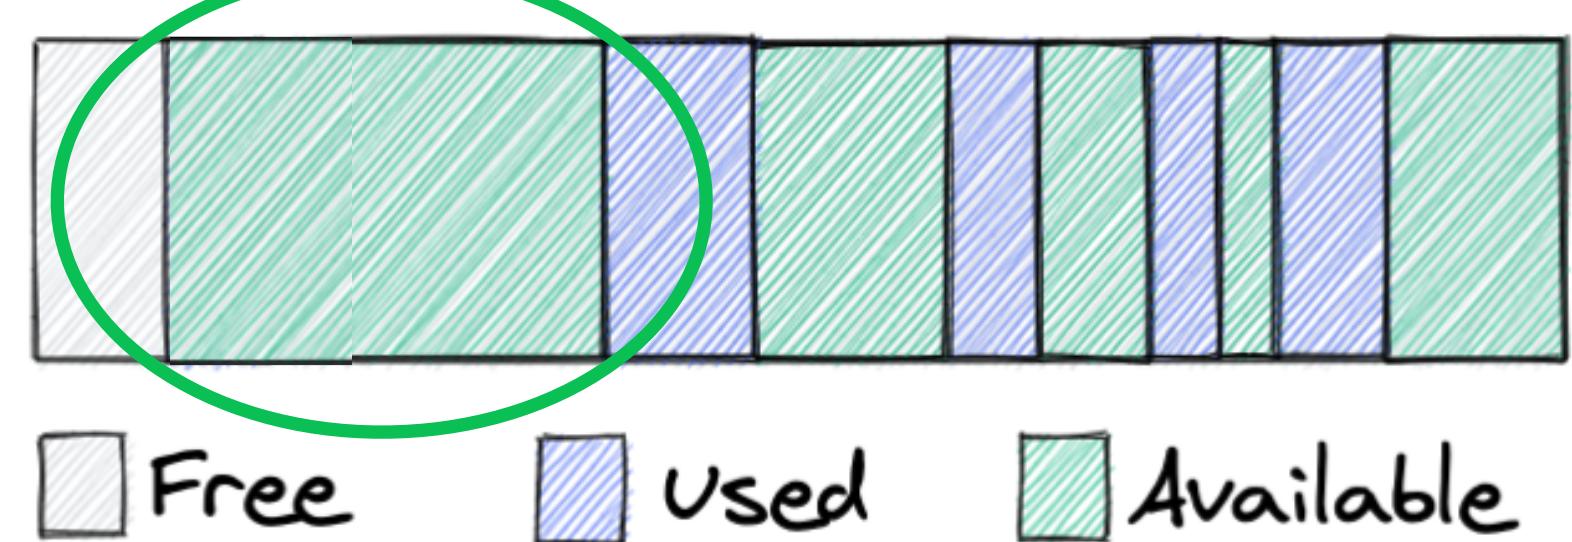
# glibc malloc

+

## Releasing memory



Will not release back



Will potentially release back

# Solving the growing memory issue

**adyen**

engineered  
for ambition

# glibc Memory Allocation Tunables

- **MALLOC\_ARENA\_MAX**
- **M\_MMAP\_THRESHOLD**
- **MALLOC\_TRIM\_THRESHOLD\_**
- **Amongst others**

# glibc Memory Allocation Tunables

- **MALLOC\_ARENA\_MAX**
  - The most used tunable, set to 2-4
  - Widely accepted as ‘the’ fix (next to switching malloc implementation)

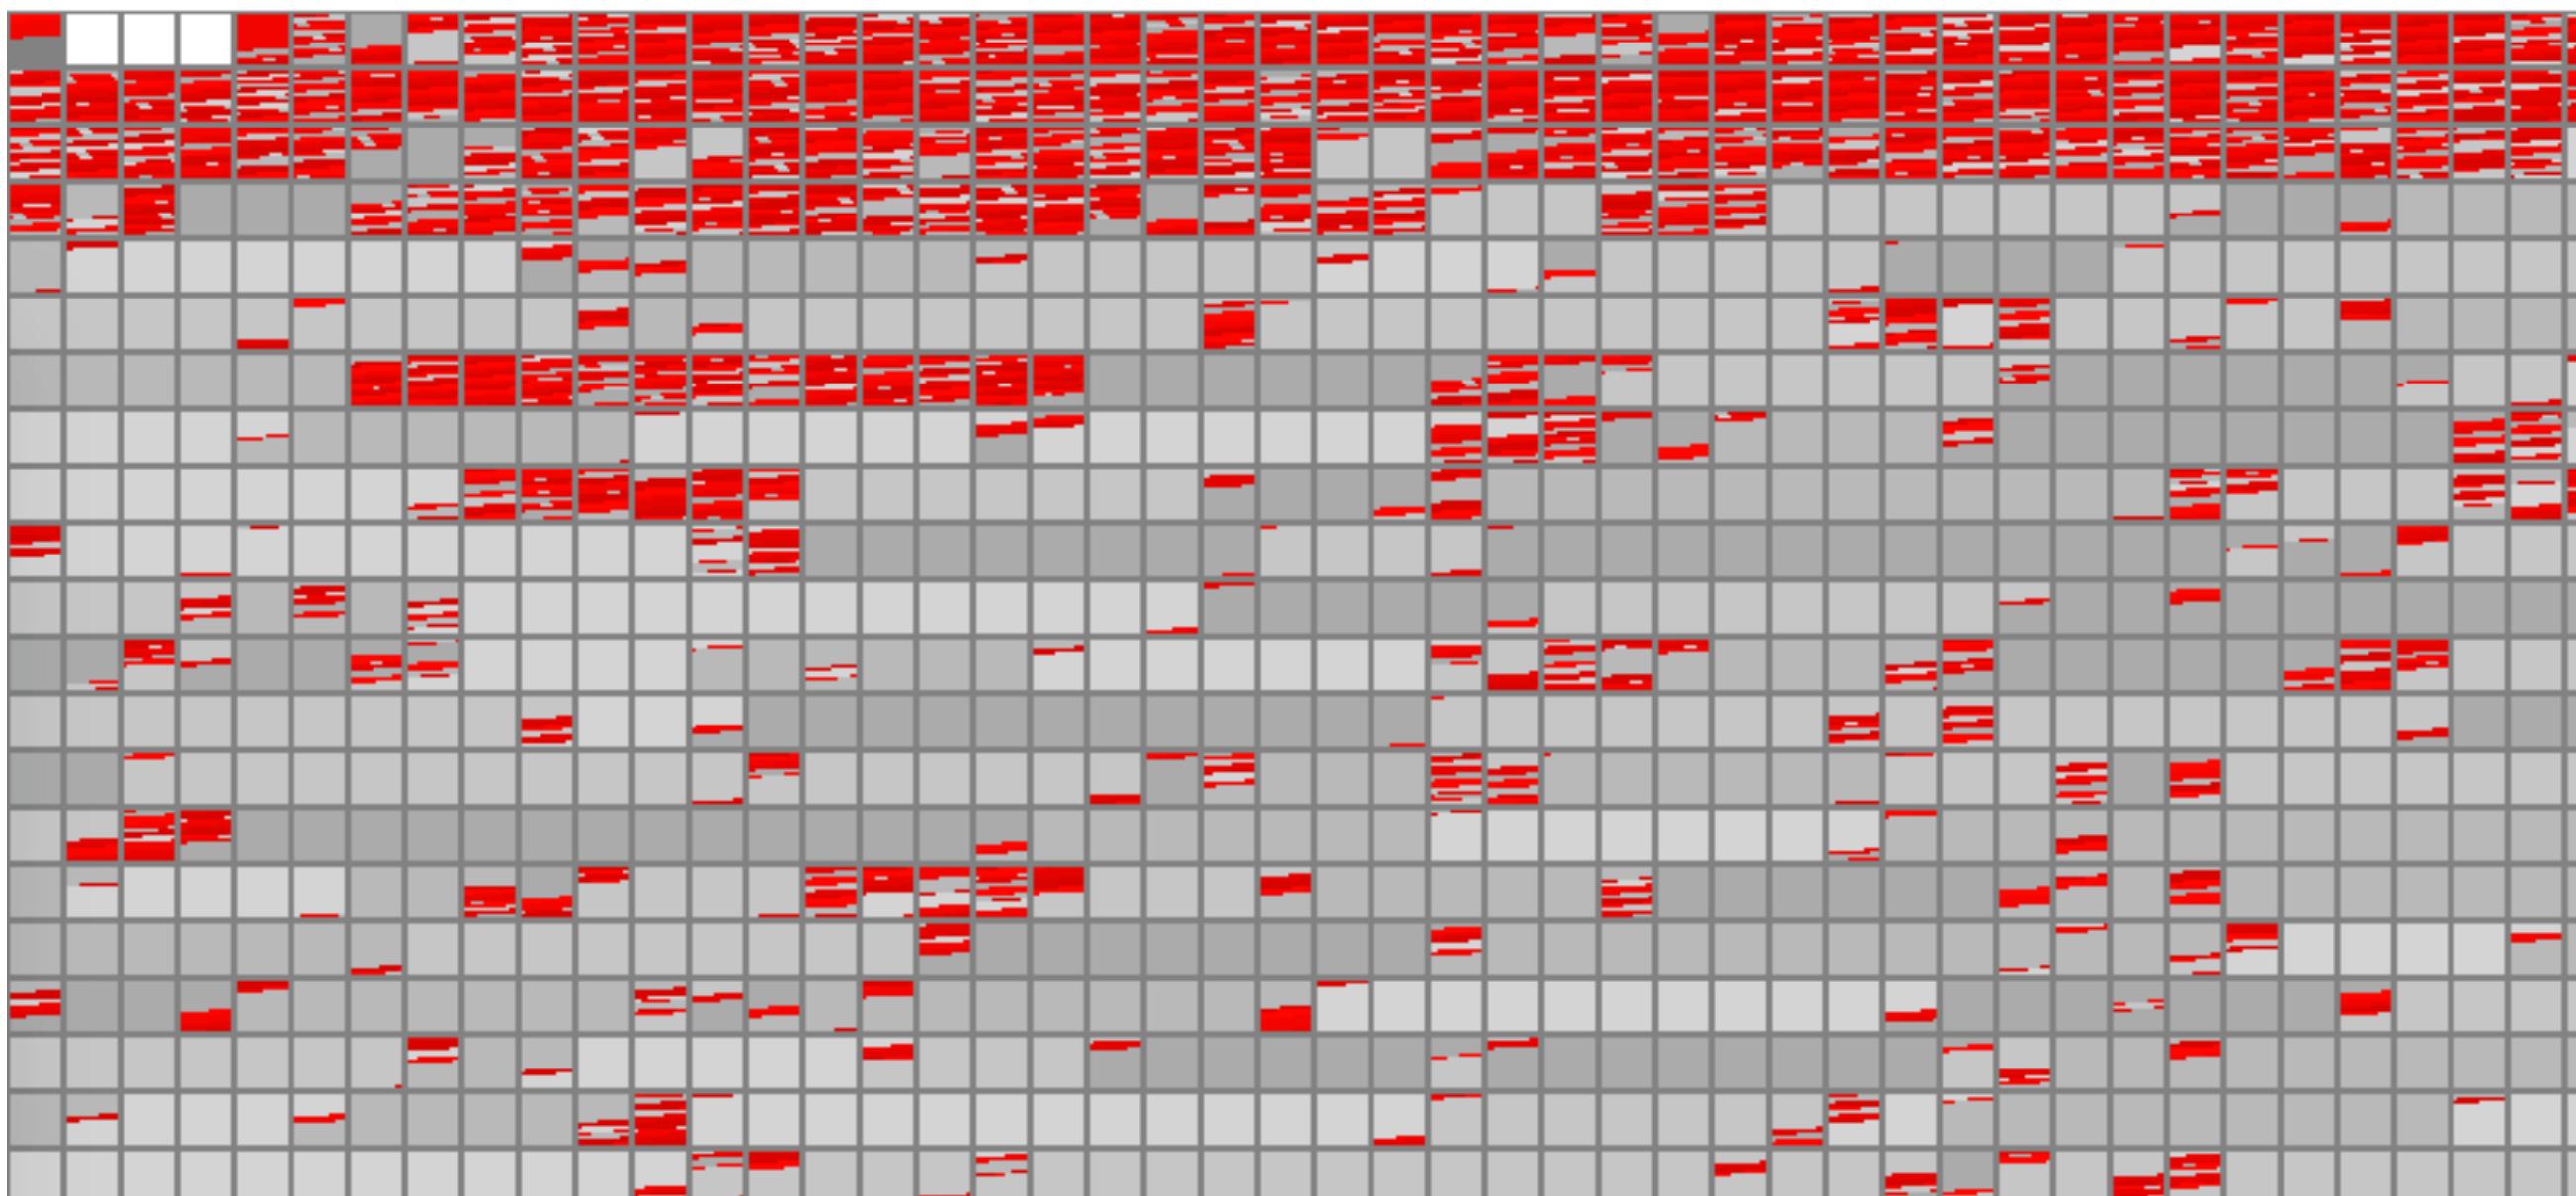
# glibc malloc\_trim

- From the manual:
  - The malloc\_trim() function attempts to release free memory at the top of the heap.
- The manual is wrong!

# glibc malloc\_trim

## Heap 29 — 0x7ff4ec000000

Virtual	9.9 MB	2533 pages	
Dirty	9.9 MB	2530 pages	99%
Clean	0.0 MB	3 pages	0%

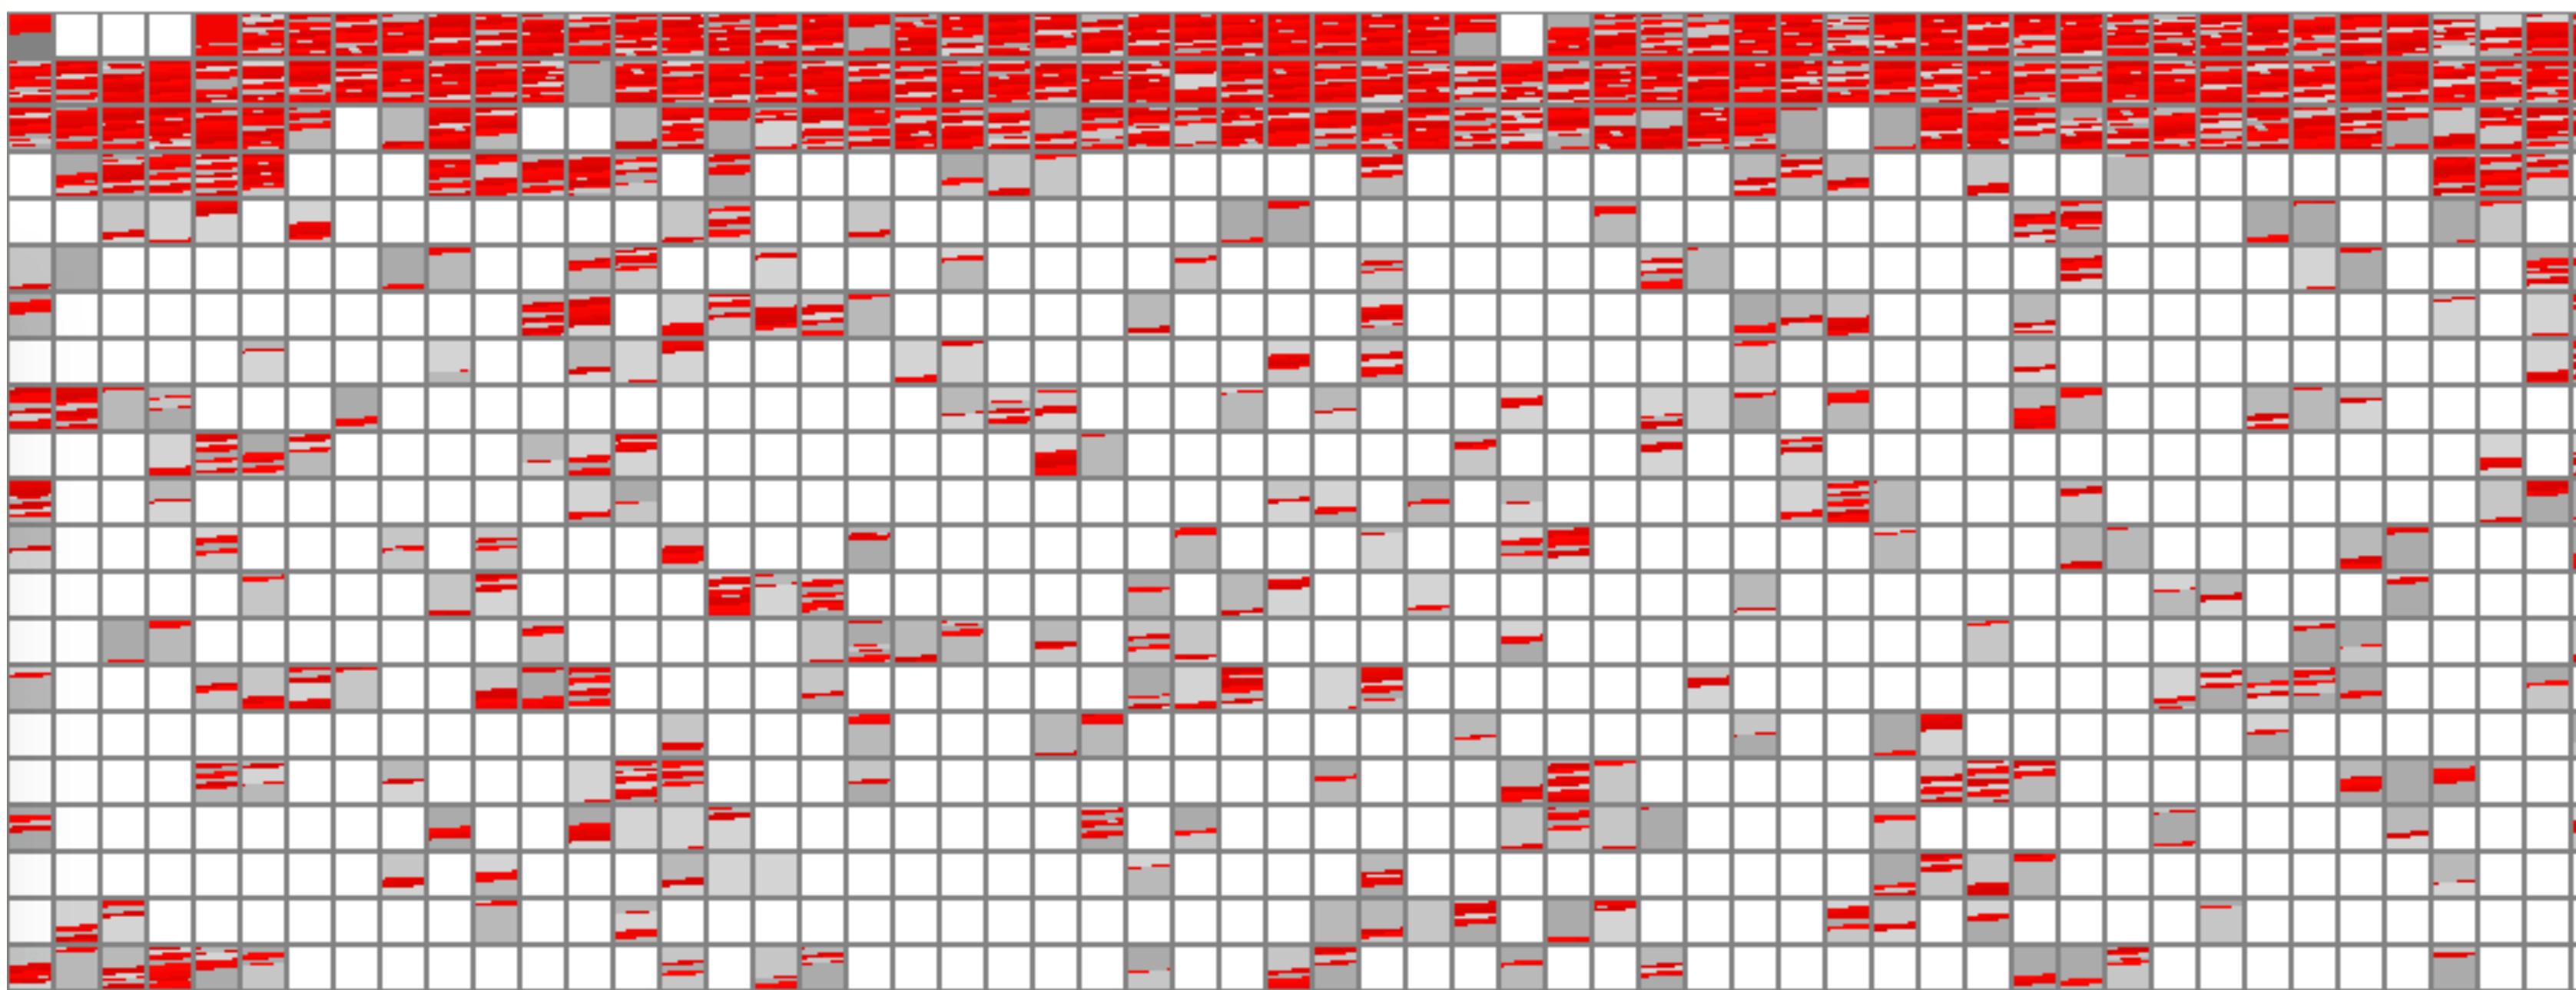


Hongli Lai - <https://www.joyfulbikeshedding.com/blog/2019-03-14-what-causes-ruby-memory-bloat.html>

# glibc malloc\_trim

## Heap 10 — 0x7f10b0000000

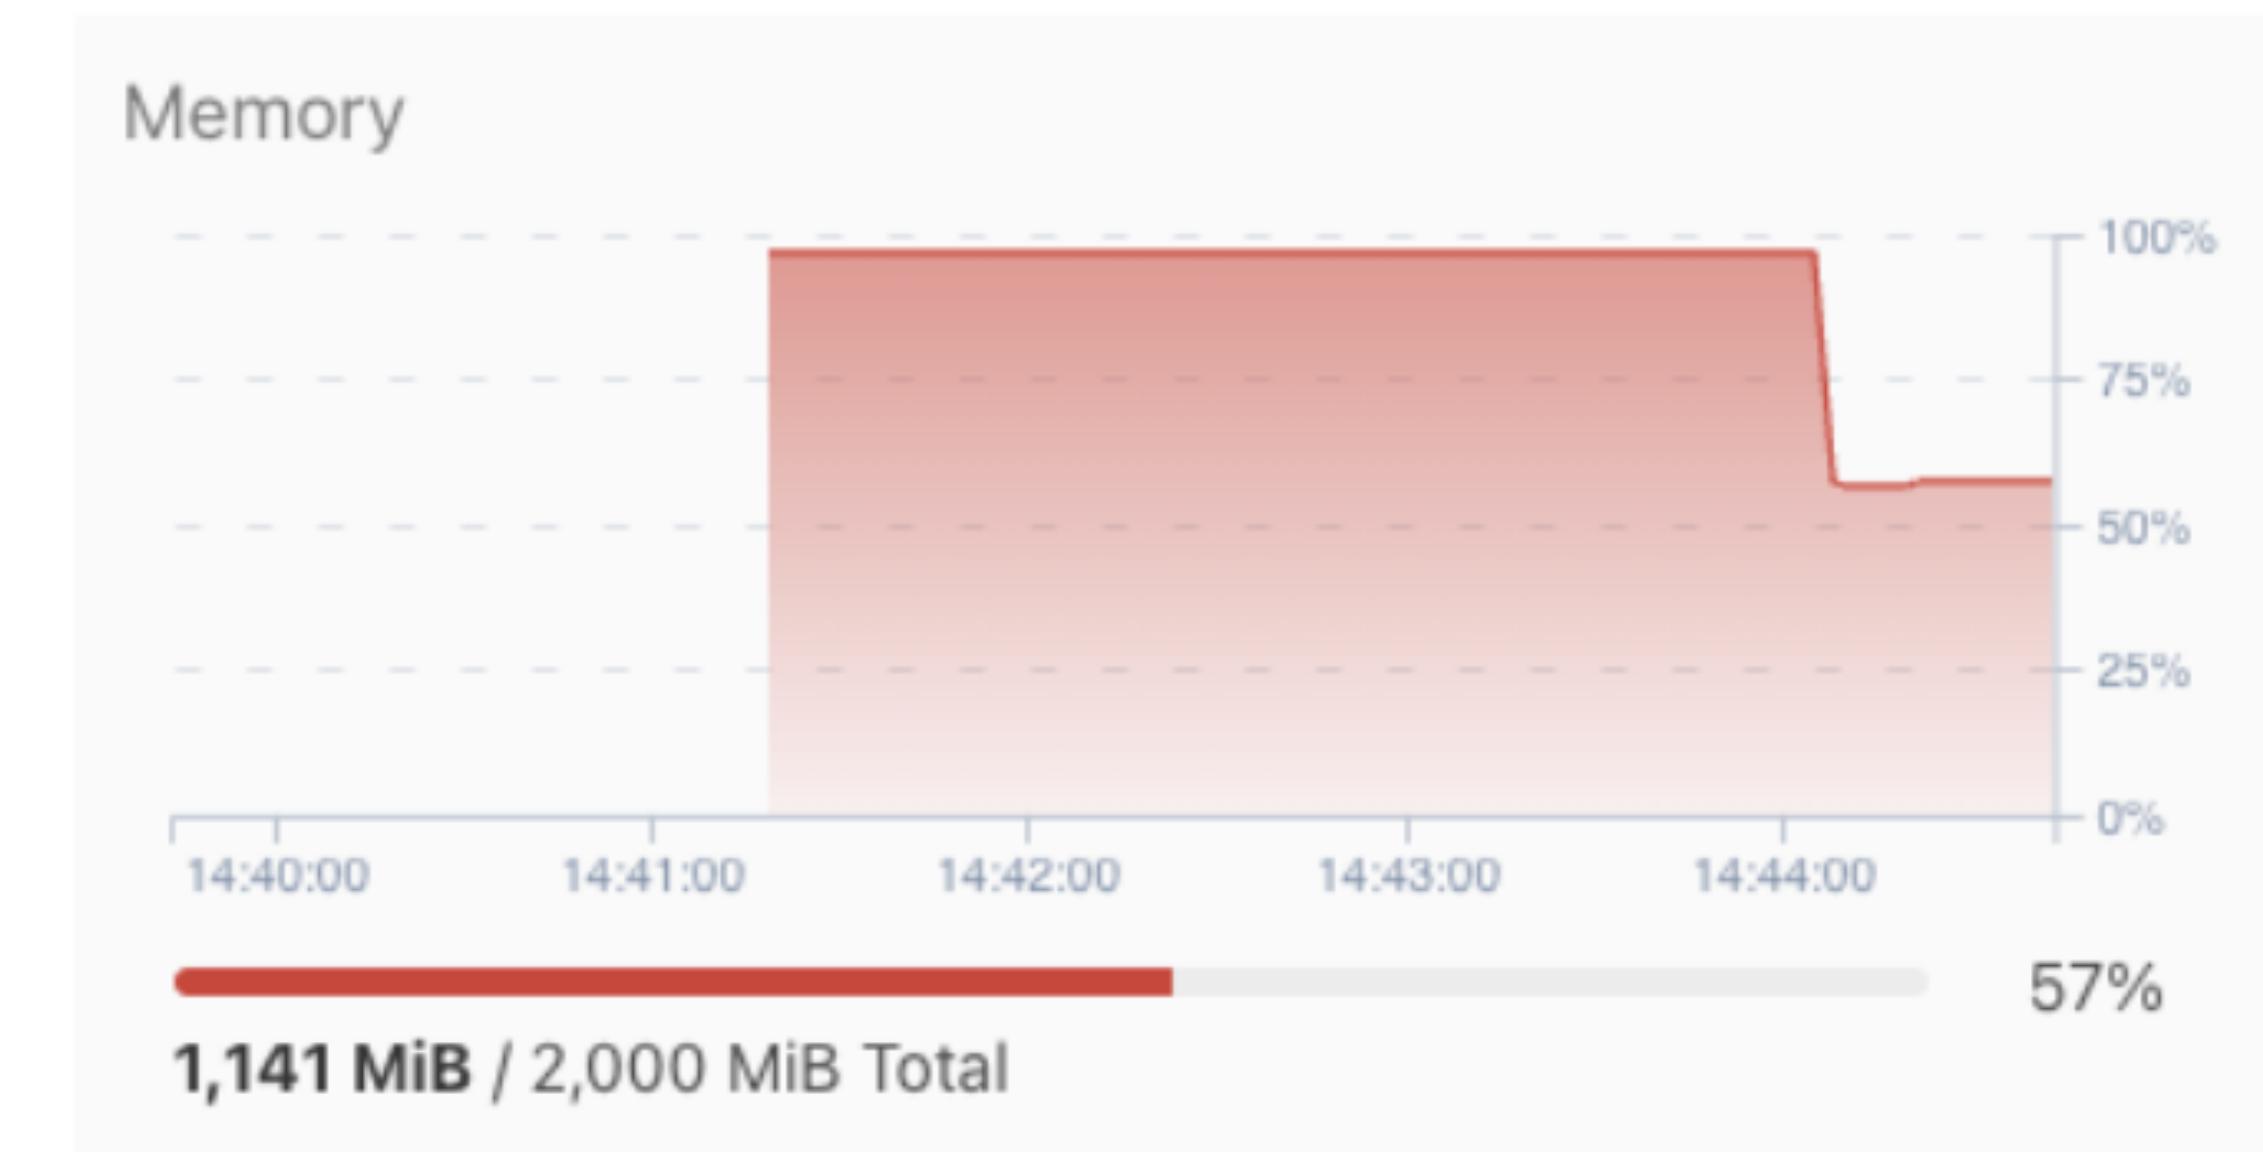
Virtual	10.7 MB	2750 pages
Dirty	2.7 MB	689 pages 25%
Clean	8.1 MB	2061 pages 74%



# glibc malloc\_trim

## System.trim\_native\_heap

```
bash-4.4$ jcmd 1 System.trim_native_heap
1:
Attempting trim...
Done.
Virtual size before: 10757268k, after: 10757140k, (-128k)
RSS before: 2009392k, after: 1191544k, (-817848k)
Swap before: 0k, after: 0k, (0k)
```



# glibc malloc\_trim

**Overhead!**

**Does a lot of system calls, which is expensive**

# Other malloc implementations

- **jemalloc**
- **mimalloc**
- **tcmalloc**
- **Many others...**
- **Your own?**
- **Using LD\_PRELOAD**

# Solving the growing memory issue

- Always test properly for your application!
- Depending on the nature, performance might be affected

# Demo time



<https://github.com/MartijnVermeulen96/MeetupMemory>



engineered  
for ambition

The end  
**Discussion time (/beers)**

# Some references:

malloc(3) Linux manual page: <https://man7.org/linux/man-pages/man3/malloc.3.html>

glibc Memory Allocation Tunables: [https://www.gnu.org/software/libc/manual/html\\_node/Memory-Allocation-Tunables.html](https://www.gnu.org/software/libc/manual/html_node/Memory-Allocation-Tunables.html)

Heroku on tuning glibc malloc, setting MALLOC\_ARENA\_MAX to 2: <https://devcenter.heroku.com/articles/tuning-glibc-memory-behavior>

Picnic blogpost on malloc: <https://blog.picnic.nl/quest-to-the-os-java-native-memory-5d3ef68ffc0a>

LinkedIn switching to Jemalloc: <https://www.linkedin.com/blog/engineering/infrastructure/taming-memory-fragmentation-in-venice-with-jemalloc>

PR on openjdk “JVM should trim the native heap”: <https://github.com/openjdk/jdk/pull/14781>

IBM mentioning setting MALLOC\_ARENA\_MAX to 4: [https://www.ibm.com/docs/en/SSFHA8\\_1.2.0/troubleshooting\\_faq/malloc\\_tuning.html](https://www.ibm.com/docs/en/SSFHA8_1.2.0/troubleshooting_faq/malloc_tuning.html)

# Some references:

Mimalloc: <https://github.com/microsoft/mimalloc>

Jemalloc: <https://github.com/jemalloc/jemalloc>

Tcmalloc: <https://github.com/google/tcmalloc>

StackOverflow for LD\_PRELOAD trick: <https://stackoverflow.com/questions/426230/what-is-the-ld-preload-trick>

Youtube video - Memory footprint of a Java process by Andrei Pangin: <https://www.youtube.com/watch?v=c755fFv1Rnk>